

Revision

COST FUNCTION, ERROR, LOSS

Cost Function

$$\textit{Squared Error} = (h(x) - y)^2$$

$$\textit{Sum of Squared Errors} = \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$$

Note: superscript is not power, but representing *ith* instance/record.


$$MSE = \frac{1}{m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$$

Or equally...

$$MSE = \frac{\sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2}{m}$$

LR with One Variable: Alternative Perspective

Size (<i>Feet</i> ²)	Price \$(× 1000)
1500	190
2250	285
2740	420
2318	300
2500	350
1250	180
...	...



**What would be the
Price for this size?**

Notations:

m = Total Number of Training Samples

x = Feature

y = Label

**$(x^{(i)}, y^{(i)})$: the i th sample in the dataset
i.e., when $i = 1$, $x^{(1)} = 2250$, $y^{(1)} = 285$**

Univariate Linear Regression

□ Also called, Linear Regression with one variable or simple linear regression

$$y = mx + b$$

$$h(x) = w_0 + w_1x$$

Parameters:


$$w_0, w_1$$

Cost Function:

$$MSE = \frac{1}{m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$$

$$J(w_0, w_1) = \frac{1}{2m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$$

Will be helpful later when differentiating. This only changes scale of minimization.



Goal: Choose w_0, w_1 such that $h(x) \approx y$ for training examples (x, y)

Or...

Goal: Choose w_0, w_1 such that $h(x) - y \approx 0$

A Simplified Case

Restrict to those lines which pass through origin.

Assume $w_0 = 0$

☐ Hypothesis

$$h(x) = w_0 + w_1 x$$

☐ Parameters

$$w_0, w_1$$

☐ Cost Function

$$J(w_0, w_1) = \frac{1}{2m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$$

☐ Goal

$$\text{Minimize}_{w_0, w_1} J(w_0, w_1)$$

☐ Hypothesis

$$h(x) = w_1 x$$

☐ Parameters

$$w_1$$

☐ Cost Function


$$J(w_1) = \frac{1}{2m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$$

☐ Goal

$$\text{Minimize}_{w_1} J(w_1)$$

Visualizing Hypothesis and Cost Function

$$h(x) = w_1 x$$



$$J(w_1) = \frac{1}{2m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$$


$$J(w_1) = \frac{1}{2m} \sum_{i=1}^m (w_1 x^{(i)} - y^{(i)})^2$$

$$J(w_1)$$

$$J(1) = \frac{1}{2(3)} ((1-1)^2 + (2-2)^2 + (3-3)^2) = 0$$

$$J(0.5) = \frac{1}{6} ((0.5-1)^2 + (1-2)^2 + (1.5-3)^2) = 0.58$$

$$J(0) = \frac{1}{6} ((0-1)^2 + (0-2)^2 + (0-3)^2) = 2.3$$



In actual, the cost functions are not that nice!

Using Both Parameters

☐ Hypothesis

$$h(x) = w_0 + w_1 x$$

☐ Parameters

$$w_0, w_1$$

☐ Cost Function

$$J(w_0, w_1) = \frac{1}{2m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$$


☐ Goal

$$\text{Minimize}_{w_0, w_1} J(w_0, w_1)$$


How would the “error surface” look like now?

Using Both Parameters

$$h(x) = w_0 + w_1 x$$




$$J(w_0, w_1) = \frac{1}{2m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$$



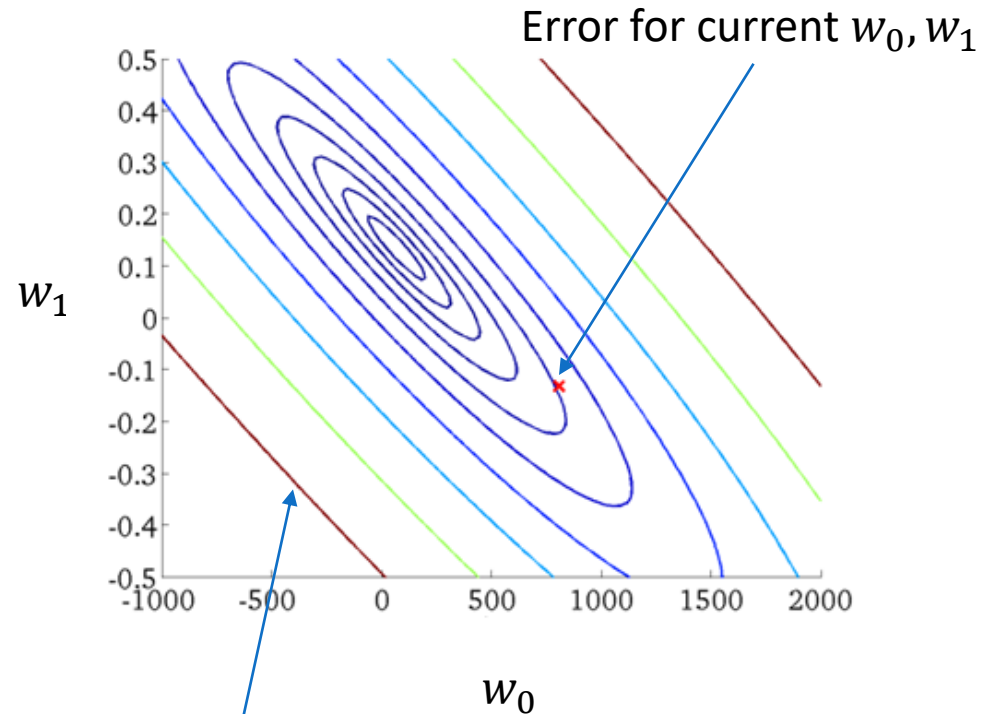
**But plotting in 3D is
not convenient!**

Contour Plots

$$h(x) = w_0 + w_1 x$$



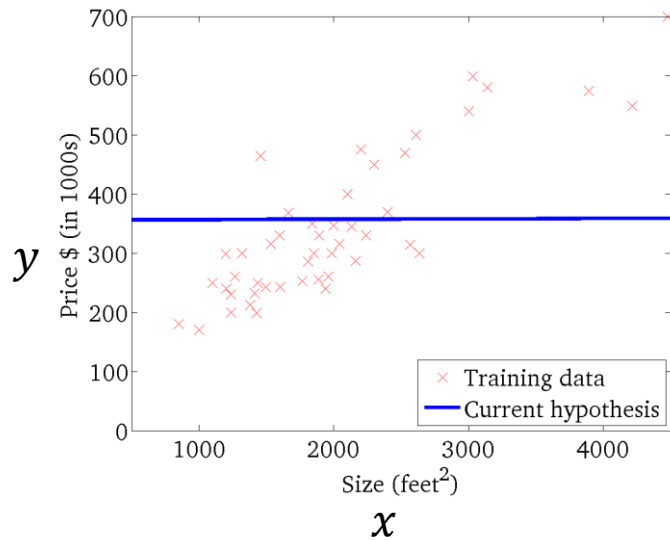
$$J(w_0, w_1) = \frac{1}{2m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$$




Highest error at edges. Lowest error as we move closer to the center.

Contour Plots

$$h(x) = w_0 + w_1 x$$




$$J(w_0, w_1) = \frac{1}{2m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$$




Contour Plots

$$h(x) = w_0 + w_1 x$$




$$J(w_0, w_1) = \frac{1}{2m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$$




Contour Plots

$$h(x) = w_0 + w_1 x$$



$$J(w_0, w_1) = \frac{1}{2m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$$



**How to find optimal
values of w_0, w_1 ?**

Gradient Descent Algorithm

The Gradient Descent Algorithm

❑ **Goal:** Minimize $J(w_0, w_1)$


❑ **Outline:**


- Start with some (w_0, w_1)
- Keep updating (w_0, w_1) to reduce $J(w_0, w_1)$
 - Until, we hopefully reach a minimum.

Global Minimum vs Local Minimum


Do we really want the Global Minimum?

We must consider test data ...





Initial values of our weights determine in which direction the algorithm would move!



A Simplified Version of Gradient Descent

- Assume again that we set $w_0 = 0$ and our hypothesis and cost function practically have only one coefficient, w_1

$$h(x) = w_1 x$$

Repeat until convergence {

$$w_1 := w_1 - \alpha \frac{\partial}{\partial w_1} (J(w_1))$$

}

Where $J(w_1) = \frac{1}{2m} \sum_{i=1}^m (w_1 x^{(i)} - y^{(i)})^2$



Step size (aka learning rate)

Direction to move

Derivate of the cost function with respect to w_1

We just want to find the rate of change from current point (instantaneous gradient)

Direction of Step




$$w_1 := w_1 - \alpha \frac{\partial}{\partial w_1} (J(w_1))$$

$w_1 := w_1 - \alpha(\text{negative}): \text{Increase } w_1$
 $w_1 := w_1 - \alpha(\text{positive}): \text{Decrease } w_1$

$$w_1 := w_1 - \alpha \frac{\partial}{\partial w_1} (J(w_1))$$

Step Size



$$w_1 := w_1 - \alpha \frac{\partial}{\partial w_1} (J(w_1))$$


$$w_1 := w_1 - \alpha \frac{\delta}{\delta w_1} (J(w_1))$$

$w_1 := w_1 - \alpha(\text{negative})$: Increase w_1

$w_1 := w_1 - \alpha(\text{positive})$: Decrease w_1

- In addition to α , the steepness of the gradient also control the step size.
 - The step sizes become smaller as we get closer to the minimum, even with a fixed α
 - With α too small, it takes a long time to reach the minimum.
-
- With α too small, it takes a long time to reach the minimum.

Step Size



$$w_1 := w_1 - \alpha \frac{\partial}{\partial w_1} (J(w_1))$$

$$w_1 := w_1 - \alpha \frac{\delta}{\delta w_1} (J(w_1))$$

$w_1 := w_1 - \alpha(\text{negative})$: Increase w_1


$w_1 := w_1 - \alpha(\text{positive})$: Decrease w_1

- In addition to α , the steepness of the gradient also control the step size.
- The step sizes become smaller as we get closer to the minimum, even with a fixed α
- With α too small, it takes a long time to reach the minimum.

- With α too small, it takes a long time to reach the minimum.

- With α too large, we can miss the minimum, and may fail to converge.

The Problem of Local Optima



Assignment 2 – Task 2

- ☐ Use your images dataset and age label (Same as Task 1 of Assignment 2)
 - Resize images to 32x32, just like before.
- ☐ Train SGD Linear Regressor using scikit-learn using training split
- ☐ Compute R^2 and MSE on test split.

- ☐ Compare the metrics with simple linear regression model trained with OLS that you trained in Task 1. (To see which one is better i.e., OLS vs SGD regression)

Book Reading

- ☐ Murphy – Chapter 1, Chapter 14
- ☐ Tom Mitchel (TM) – Chapter 4