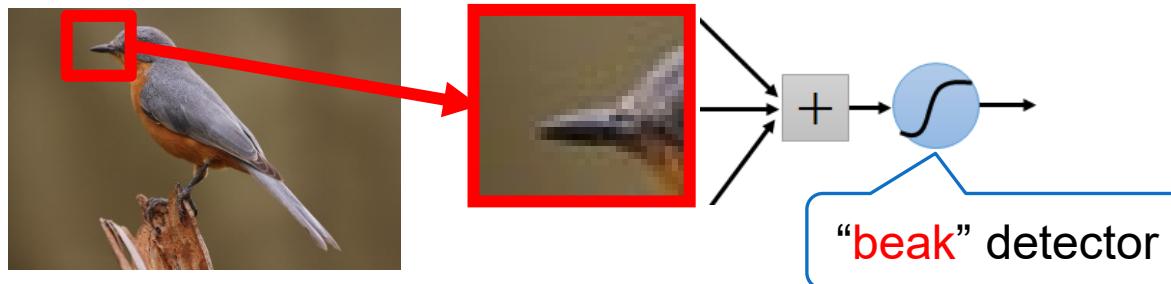


Color Images, Multiple Filters, and Parameters Sharing

Consider Learning an Image

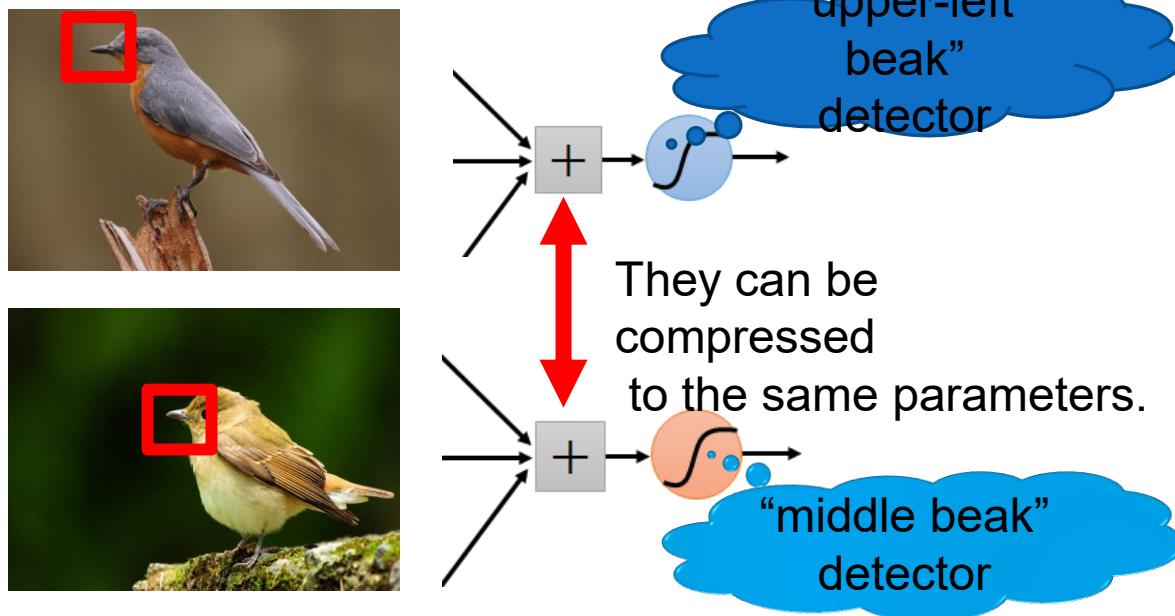
- ❑ Some patterns are much smaller than the whole image

Can represent a small region with fewer parameters



What if Pattern Location Changes?

- ❑ What about training a lot of such “small” detectors, and each detector must “move around”?



Convolution

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

These are the network parameters to be learned.

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

: :

Each filter detects a small pattern (3 x 3).

Convolution

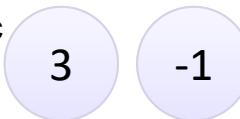
1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

stride=

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

Dot
product



6 x 6 image

Convolution

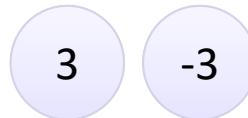
1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

If

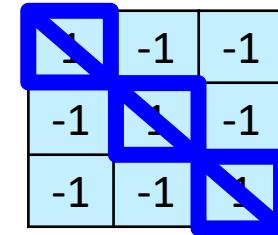
stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0



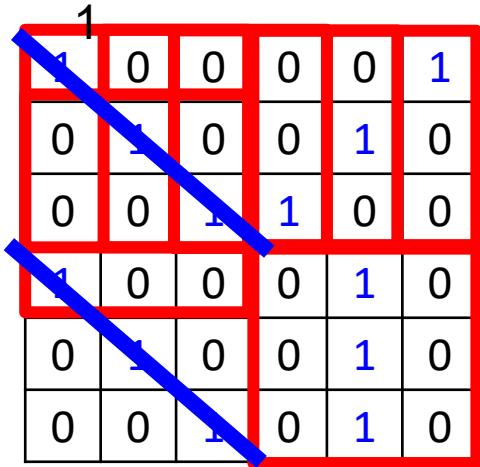
6 x 6 image

Convolution

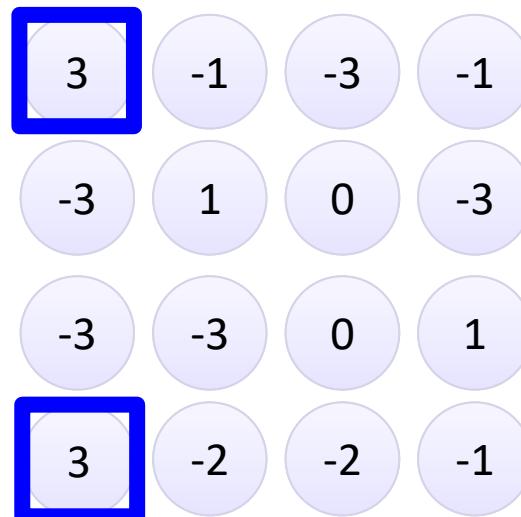


Filter 1

stride=



6 x 6 image



Convolution

-1	1	-1
-1	1	-1
-1	1	-1

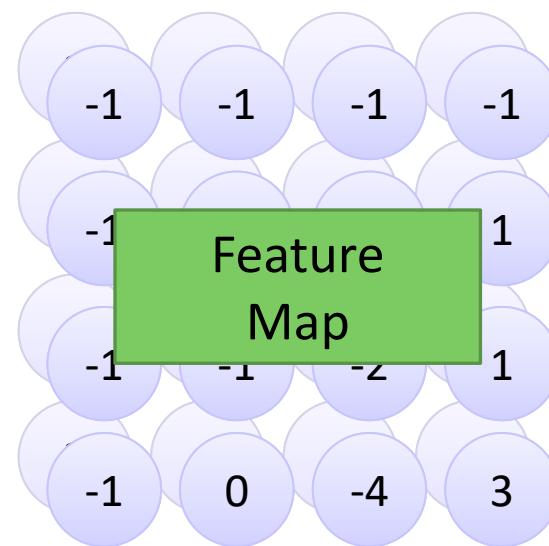
Filter 2

stride=

1	0	0	0	0	0	1
0	1	0	0	0	1	0
0	0	1	1	0	0	0
1	0	0	0	1	0	0
0	1	0	0	0	1	0
0	0	1	0	1	0	0

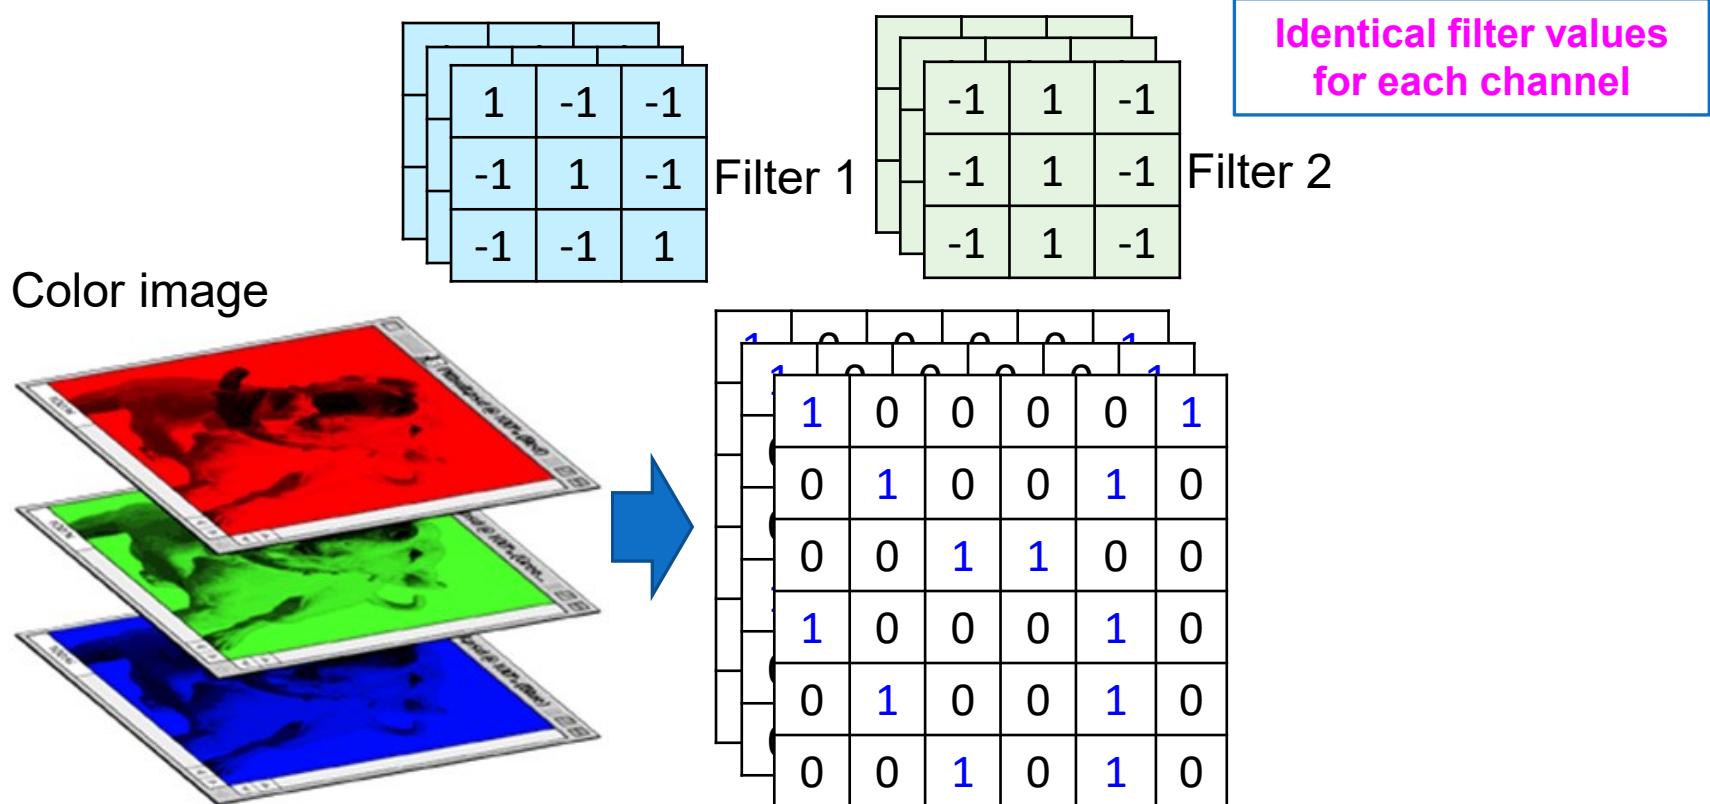
6 x 6 image

Repeat this for each filter

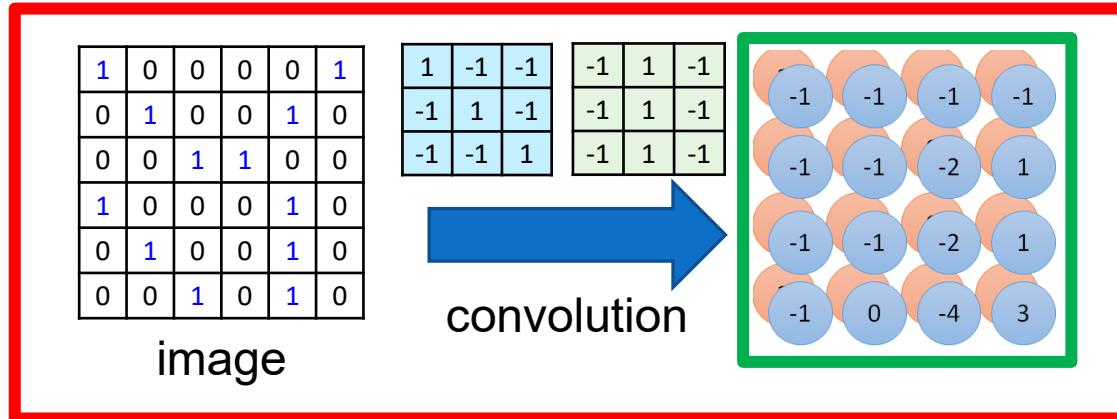


Two 4 x 4 images
Forming 4 x 4 x 2 matrix

Convolutions: Color Image

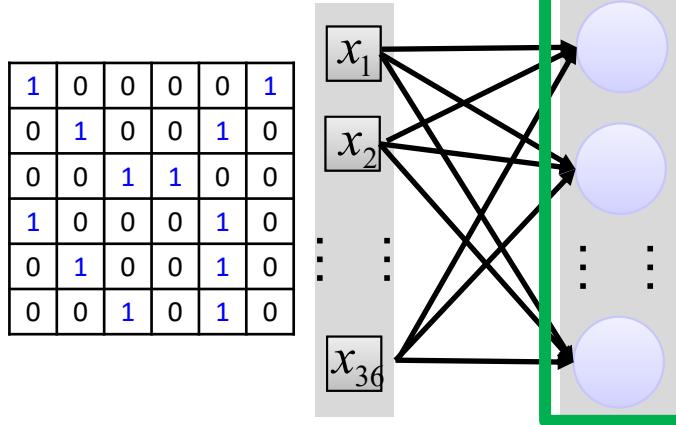


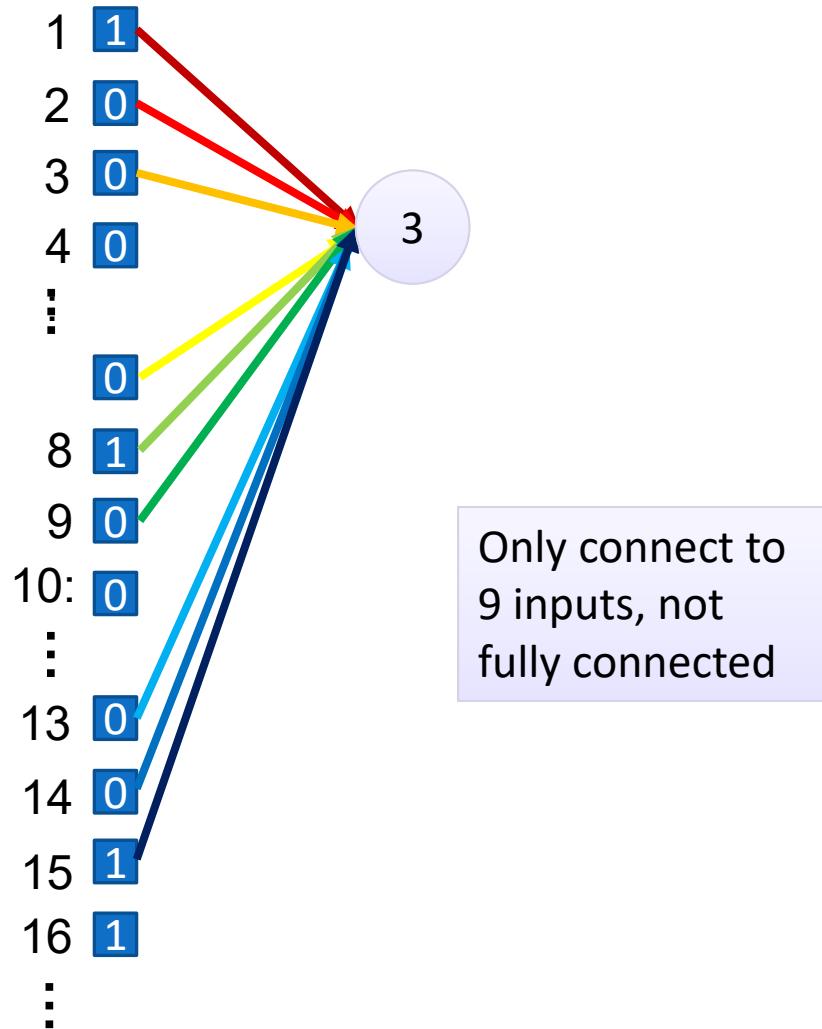
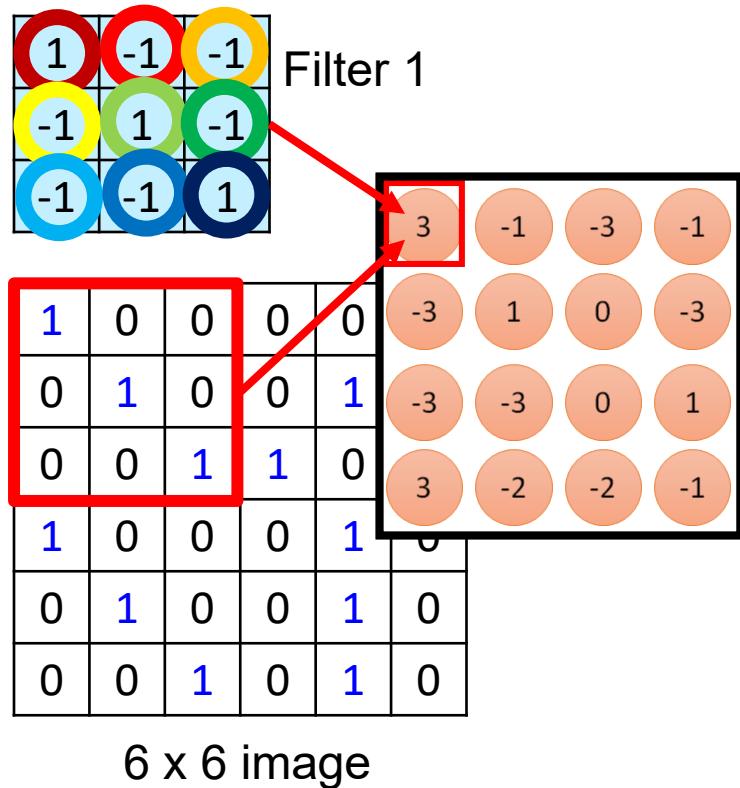
Convolutions VS Fully Connected

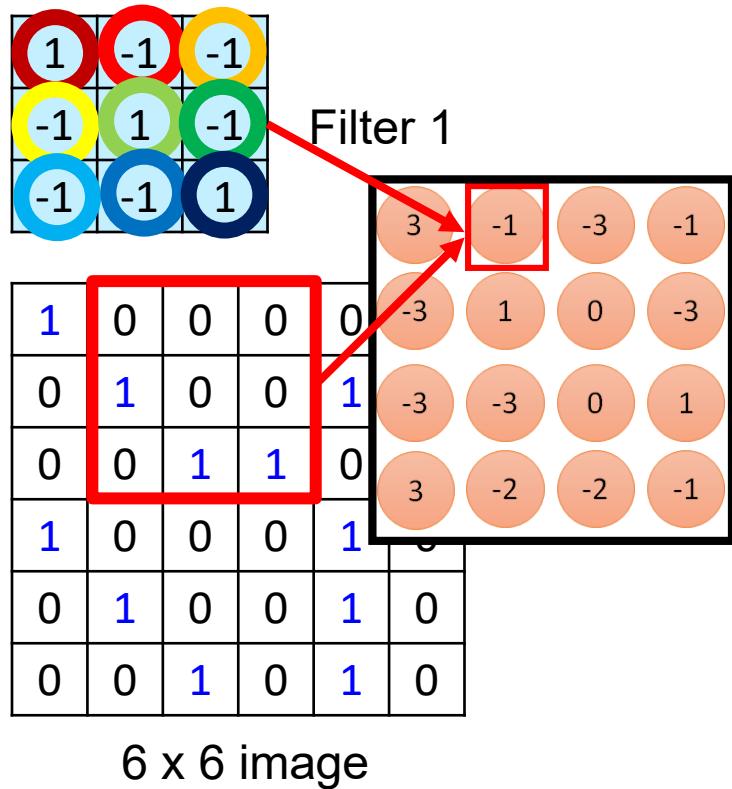


Fully-
connected

1	0	0	0	0	0	1
0	1	0	0	0	1	0
0	0	1	1	0	0	
1	0	0	0	0	1	0
0	1	0	0	0	1	0
0	0	1	0	0	1	0

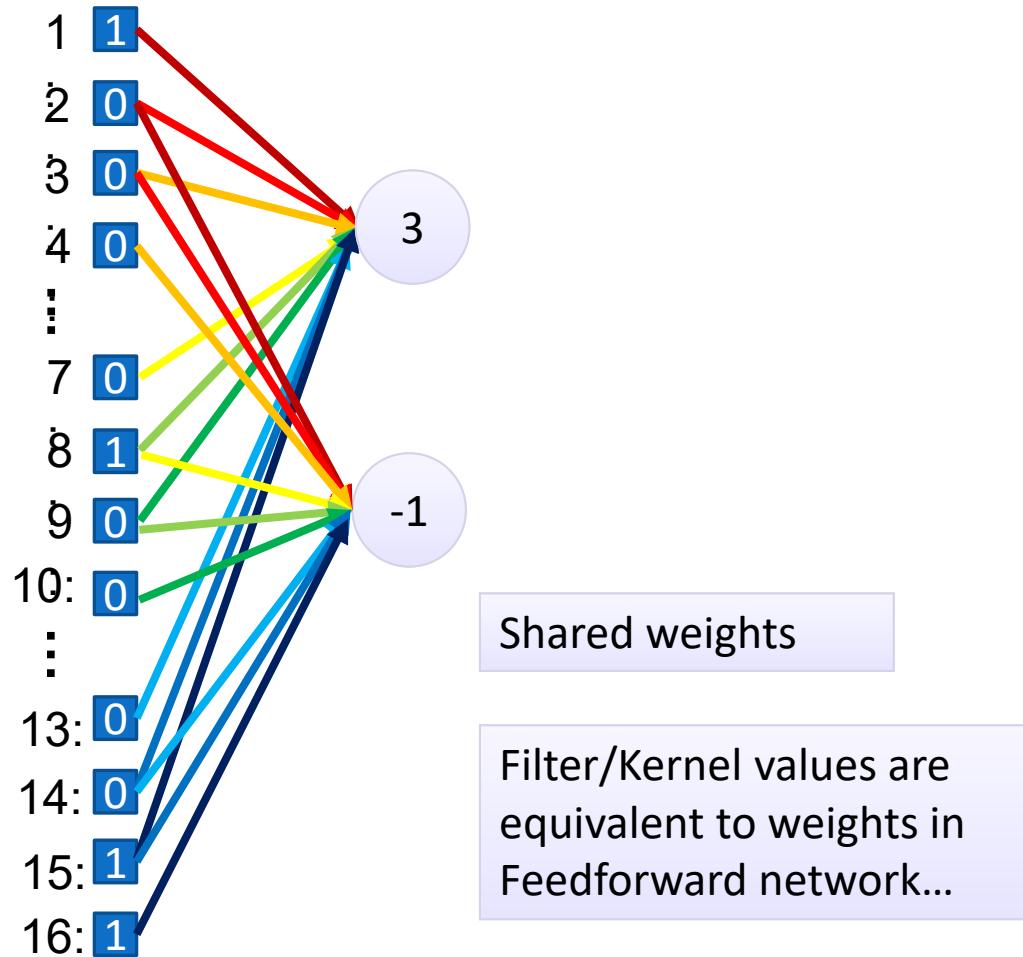






Fewer parameters

Even fewer parameters



Shared weights

Filter/Kernel values are equivalent to weights in Feedforward network...

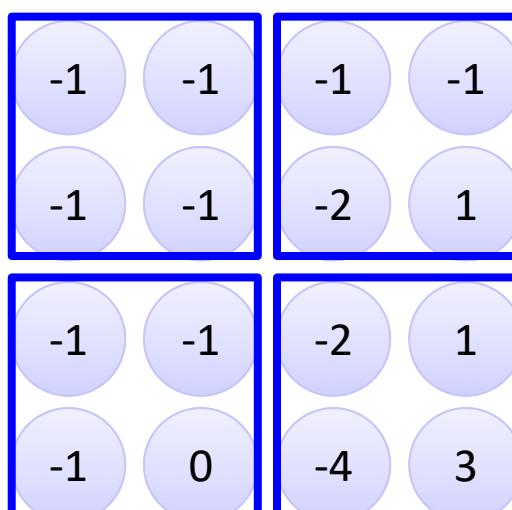
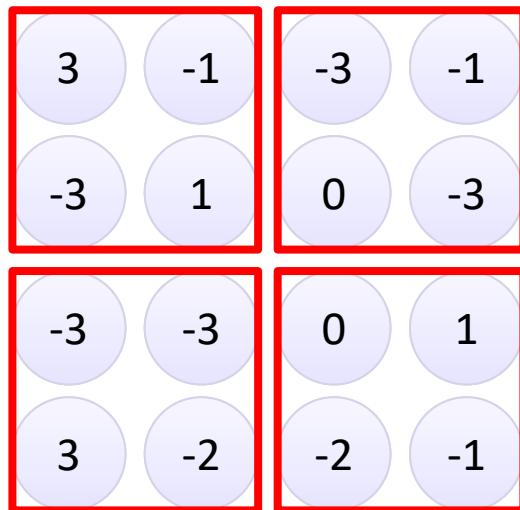
Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

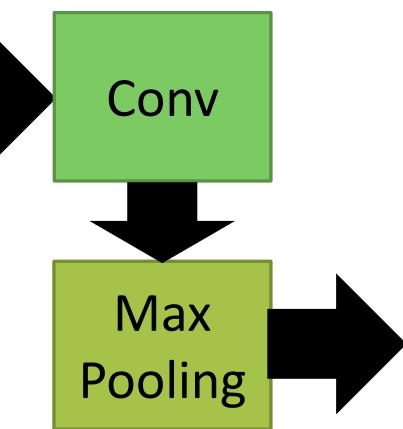
Filter 2



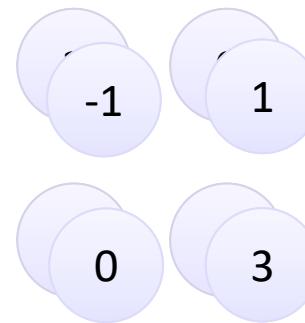
Max Pooling

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image



New image
but smaller

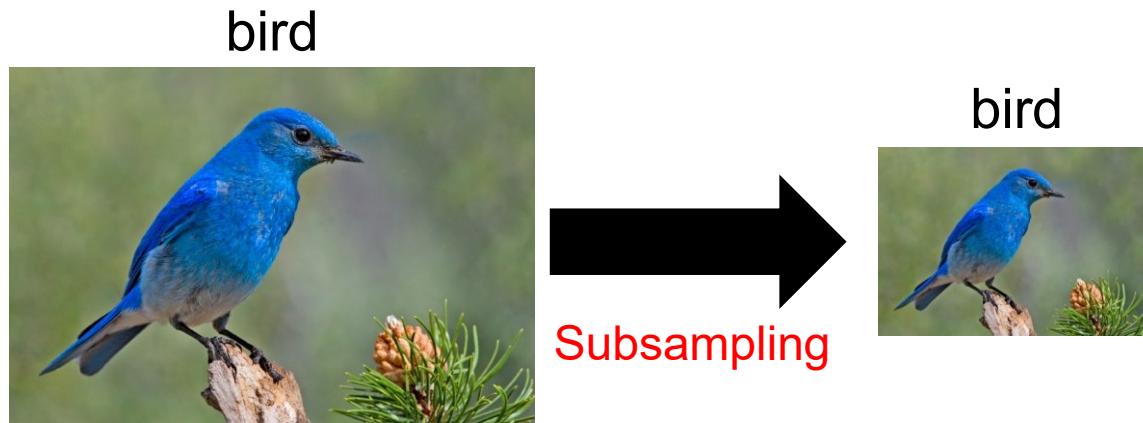


2 x 2 image

Each filter
is a channel

Why Pooling?

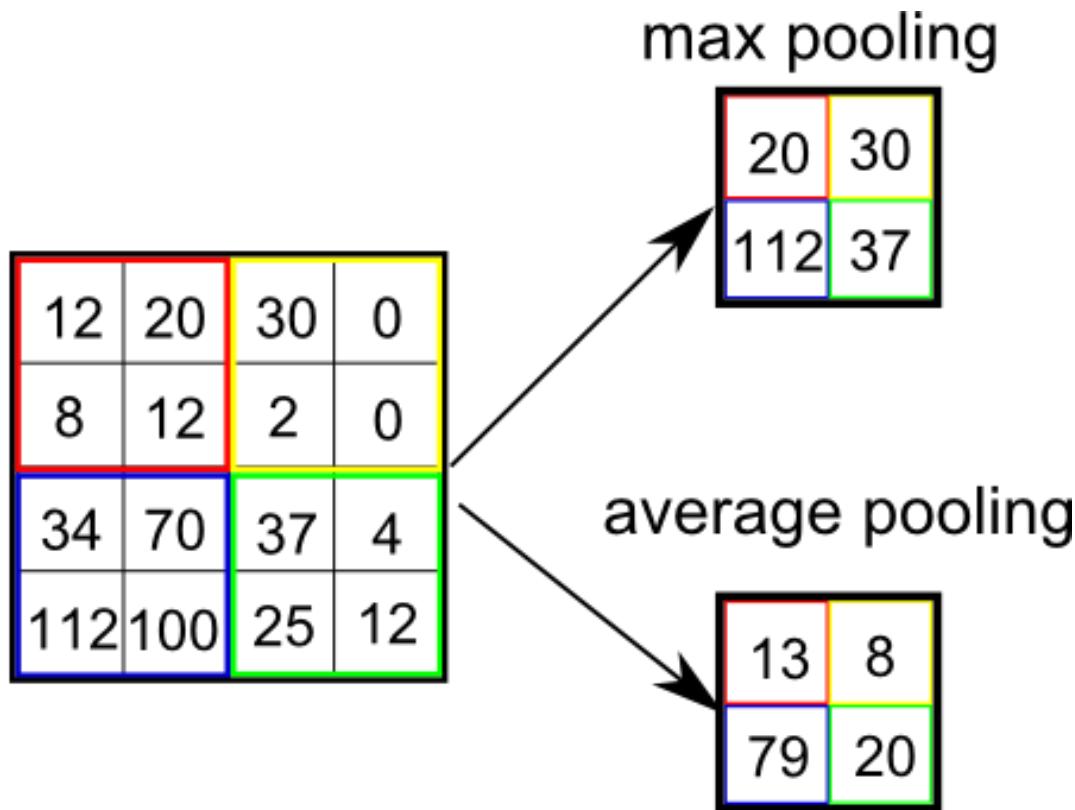
- ❑ Subsampling the pixels will **not change the object!!**



We can subsample the pixels to make image smaller

→ fewer parameters to characterize the image

Pooling Strategies

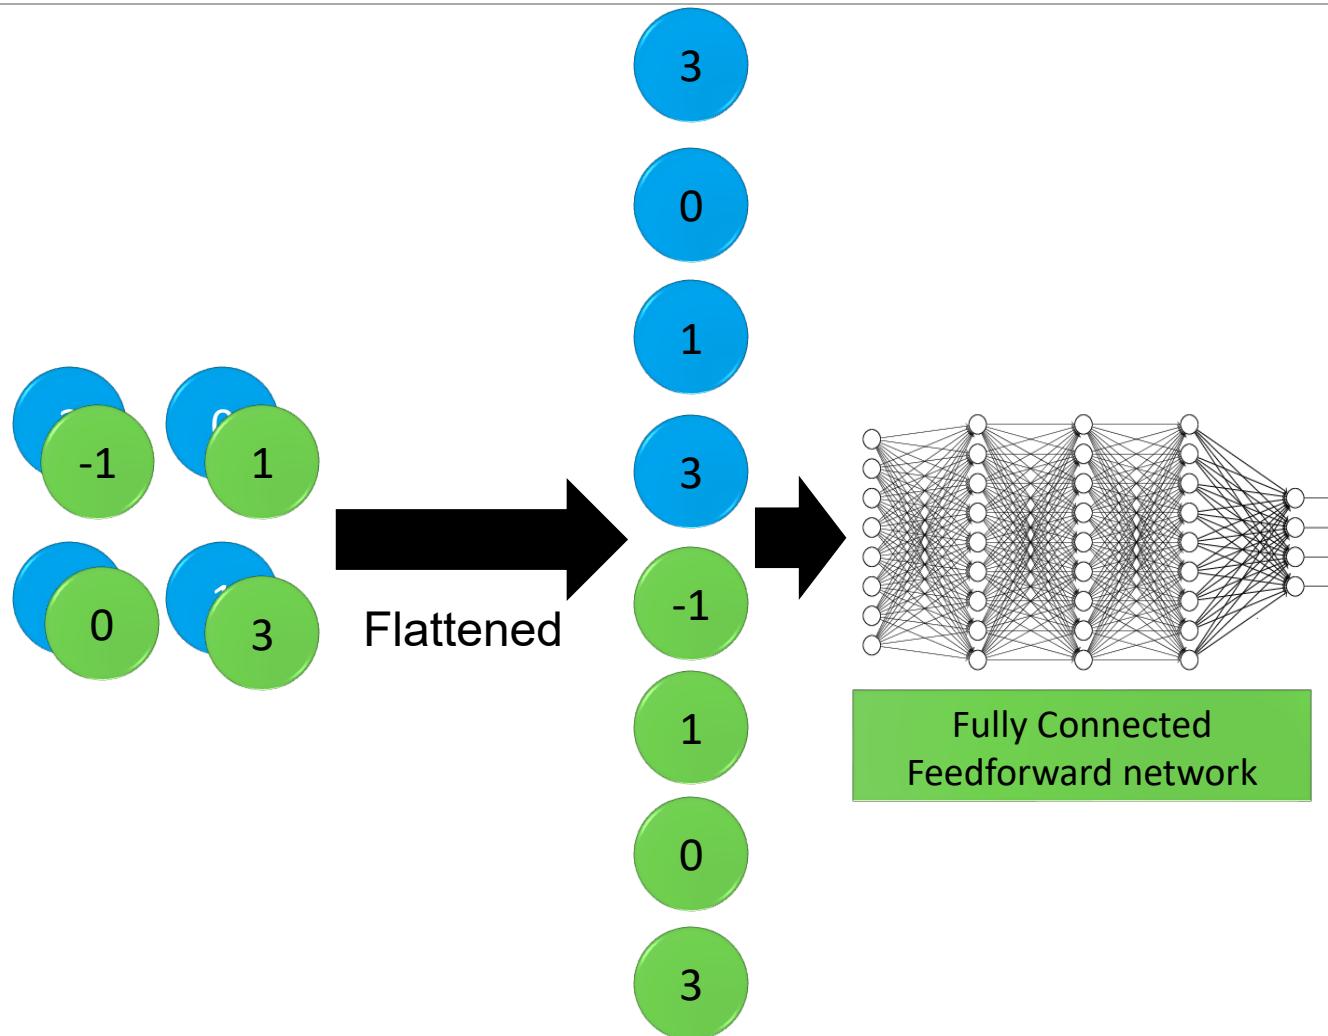


A CNN Compresses A Fully Connected Network...

- ❑ CNN Compresses a fully connected network in two ways:
 - Reducing number of connections
 - Shared weights

- ❑ Max/Average Pooling further reduces the complexity

Flattening

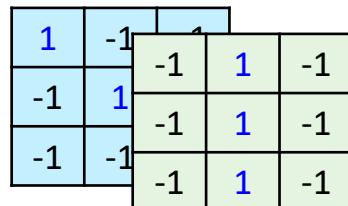


Implementation in TensorFlow Keras

USING SEQUENTIAL MODEL (CAN ALSO BE
ADAPTED TO FUNCTIONAL API)

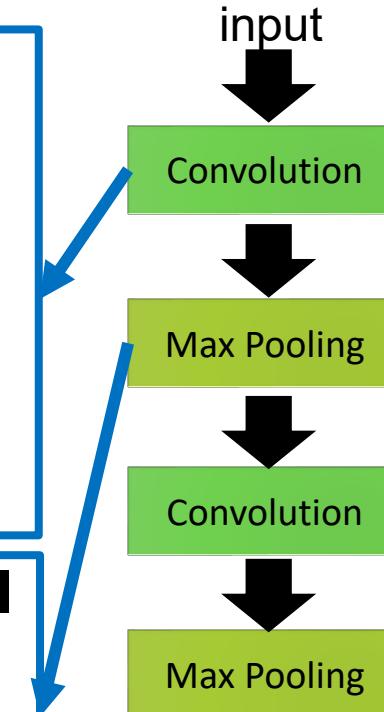
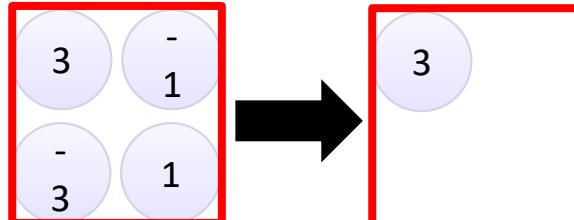
Only modified the ***network structure*** and ***input format (vector -> 3-D tensor)***

```
model2.add( Convolution2D( 25, 3, 3,  
                           input_shape=(28, 28, 1)) )
```

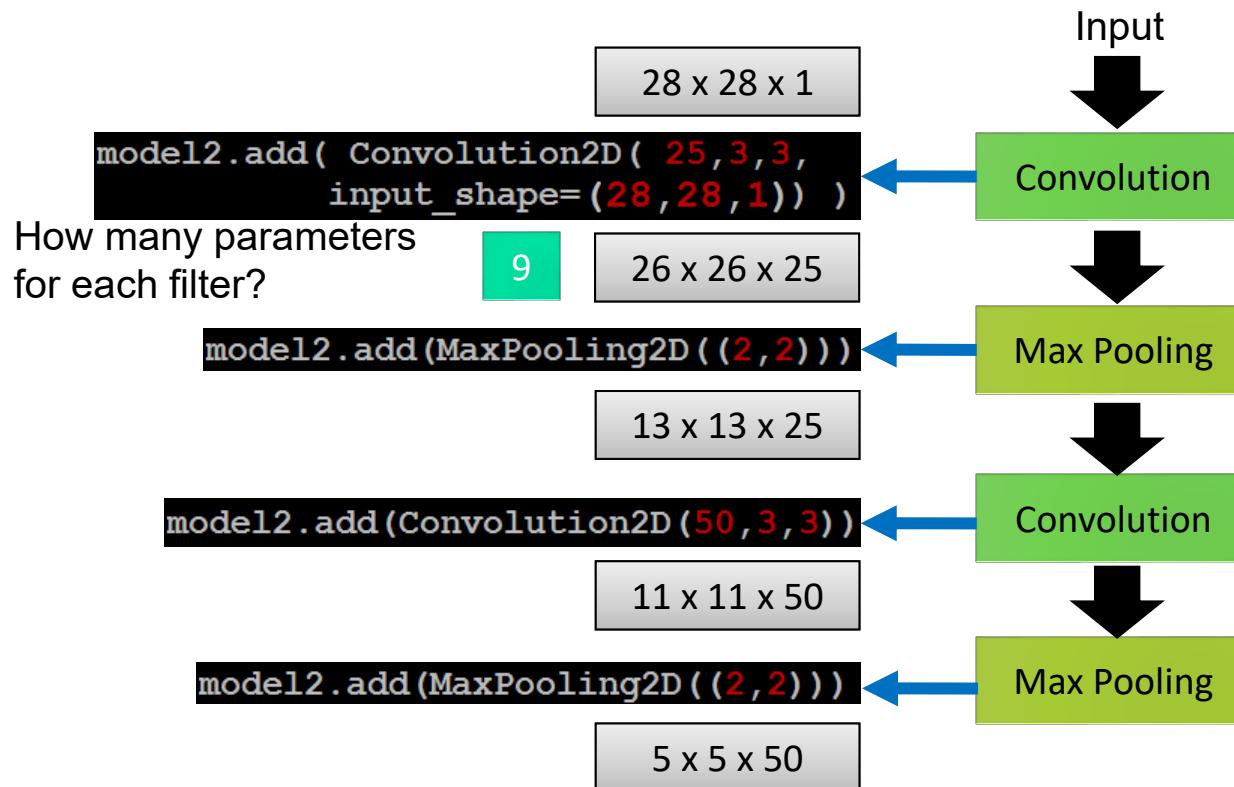


...
There are
25 3x3
filters.
...
Input_shape = (28 , 28 , 1)
28 x 28 pixels 1: black/white, 3: RGB

```
model2.add(MaxPooling2D((2, 2)))
```

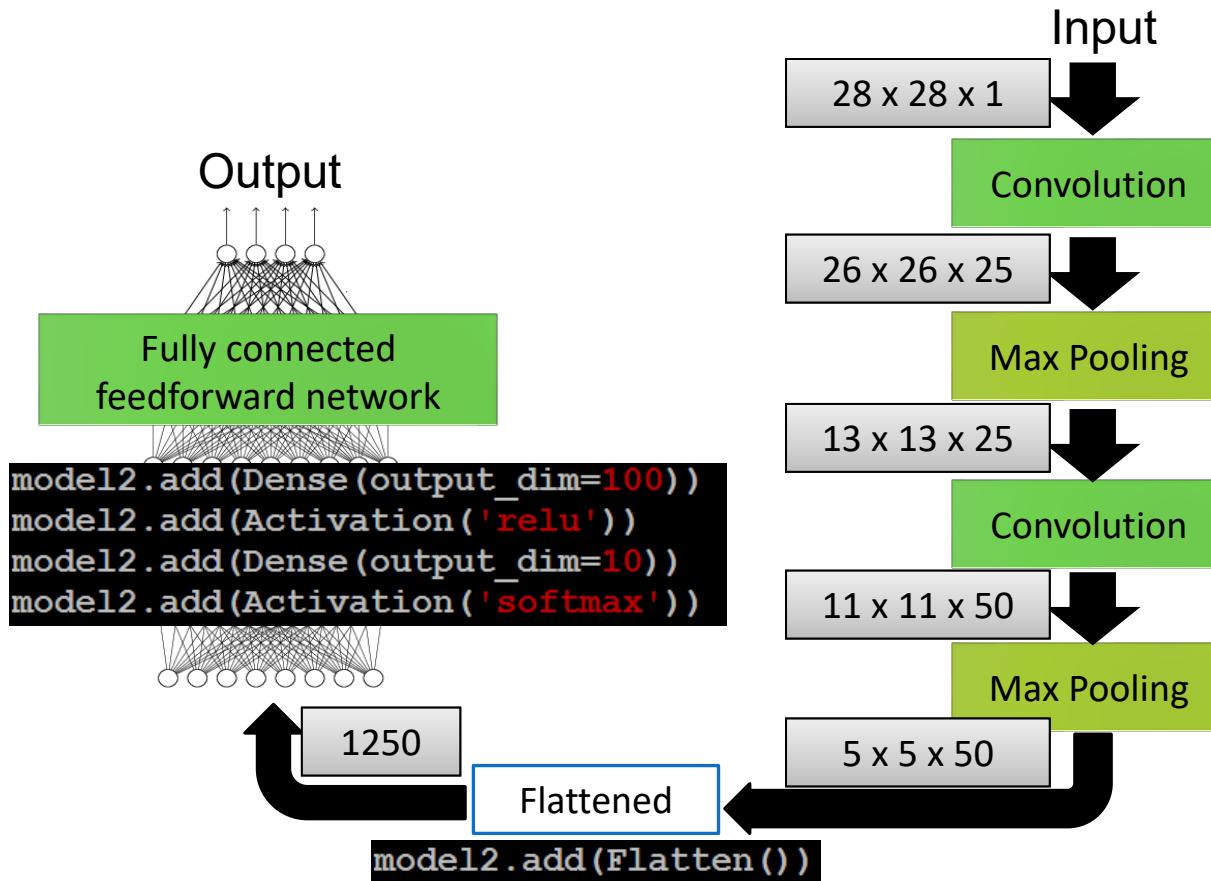


Only modified the ***network structure*** and ***input format (vector -> 3-D array)***



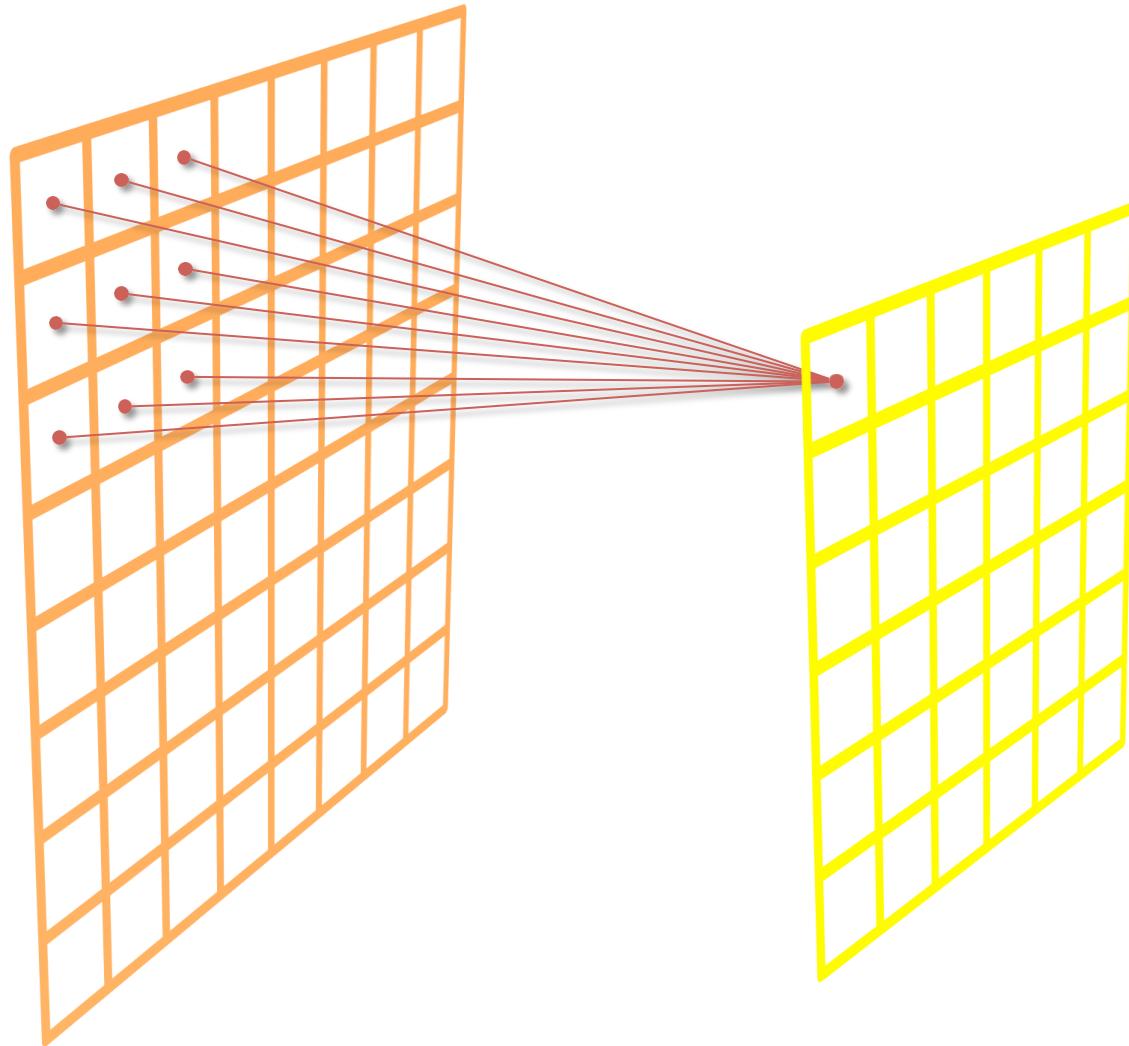
How many parameters
for each filter?

Only modified the ***network structure*** and ***input format (vector -> 3-D array)***

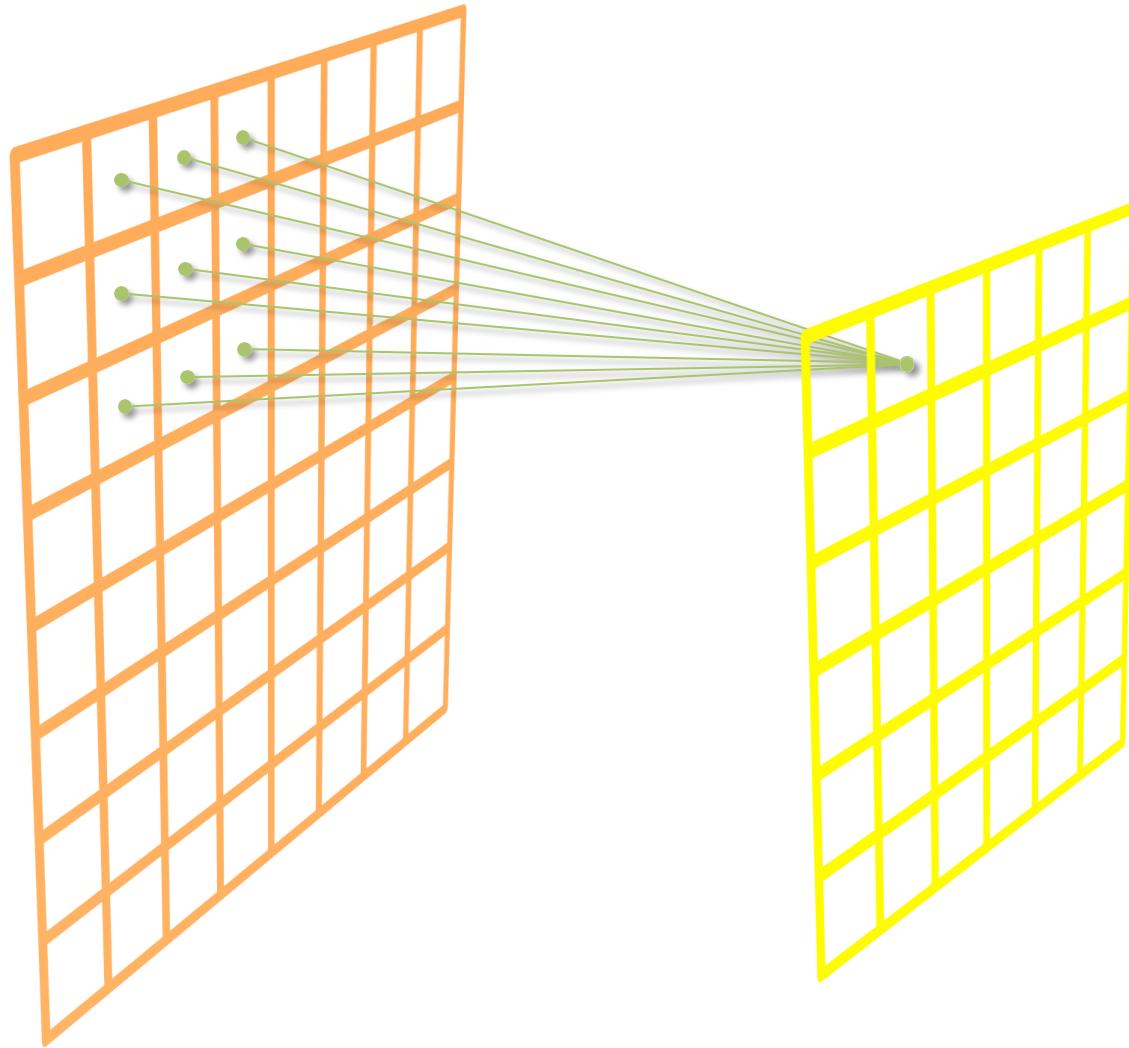


Weights Sharing in CNNs

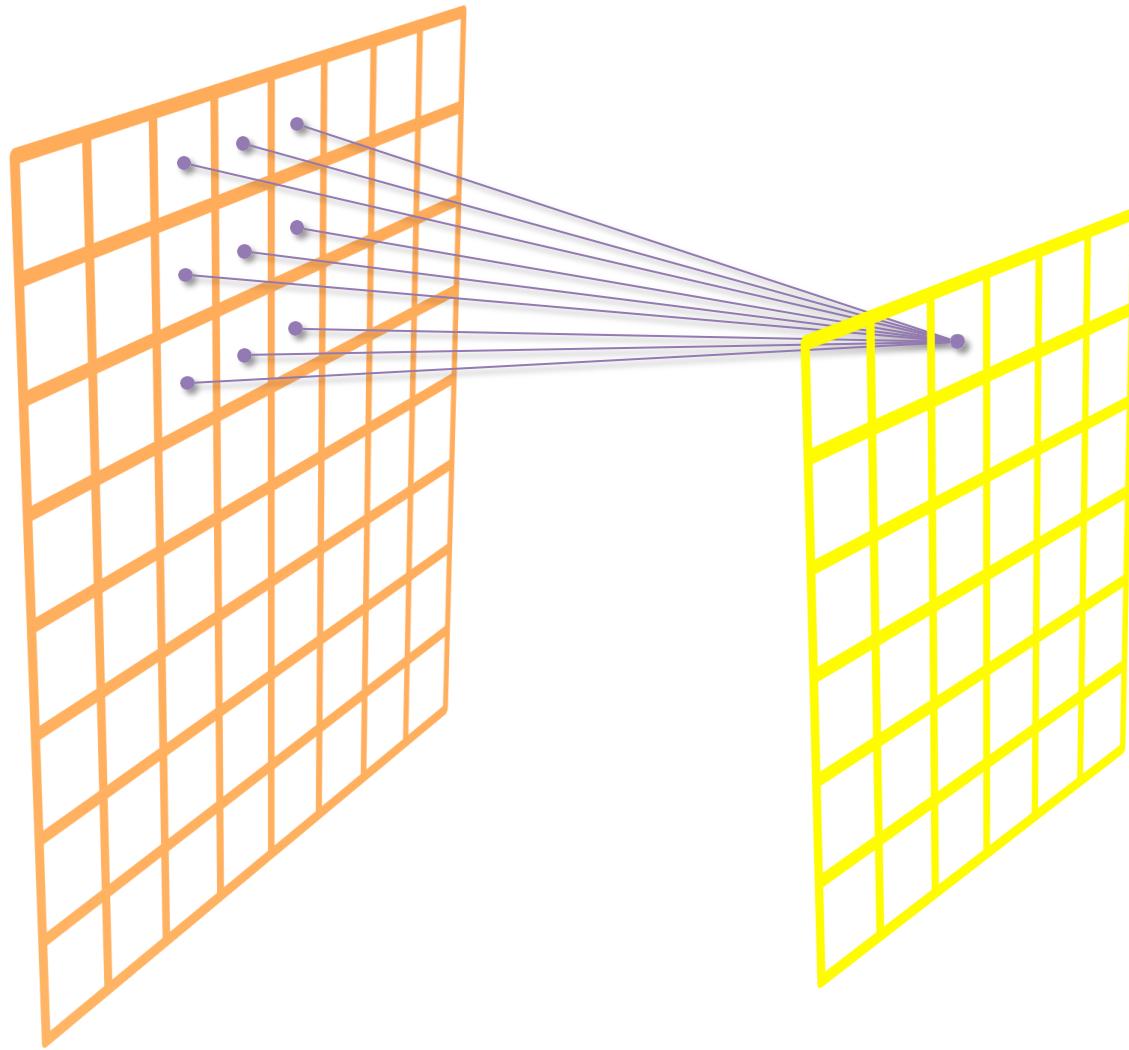
8x8 image, 3x3 filter, Stride 1



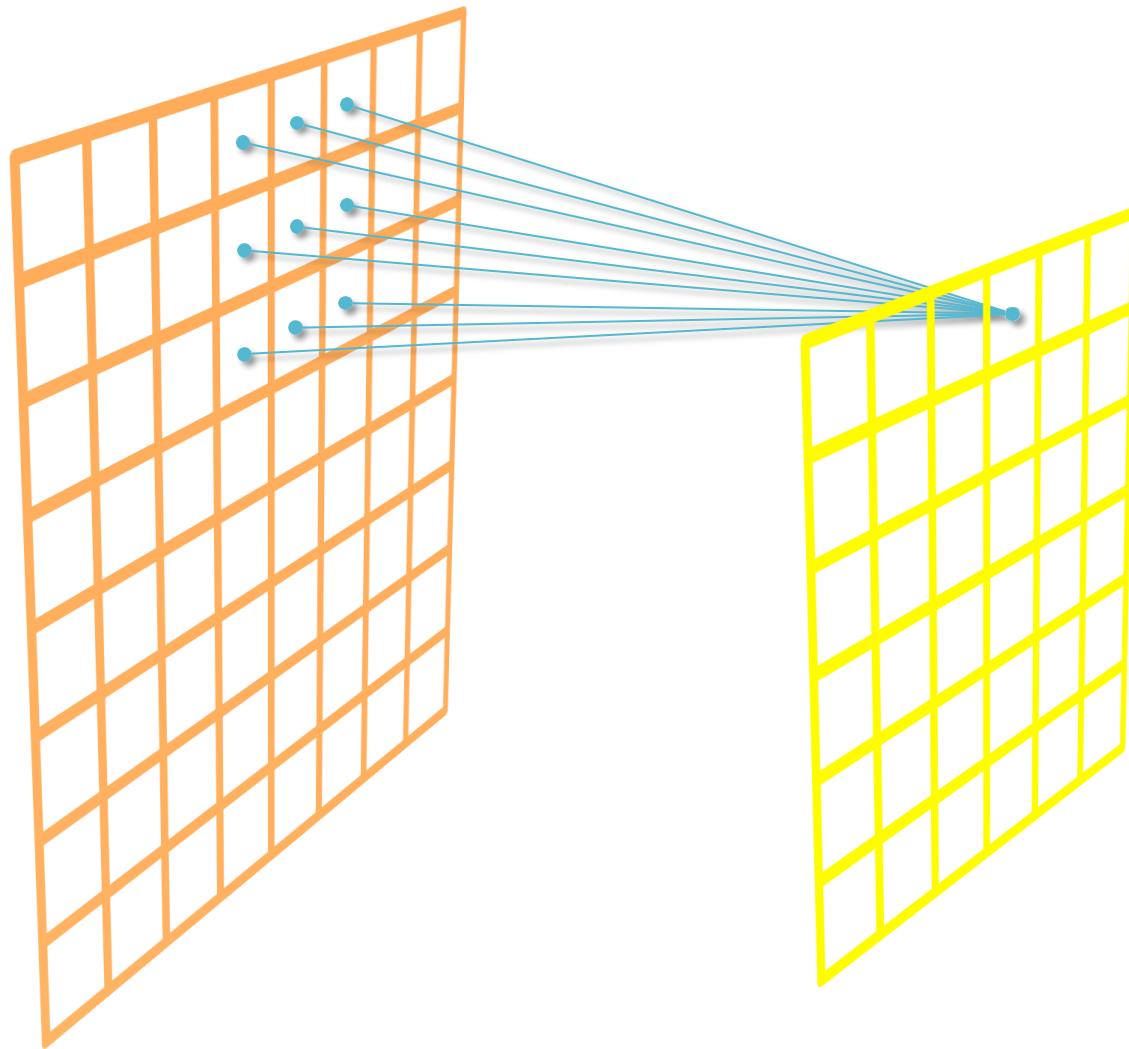
8x8 image, 3x3 filter, Stride 1



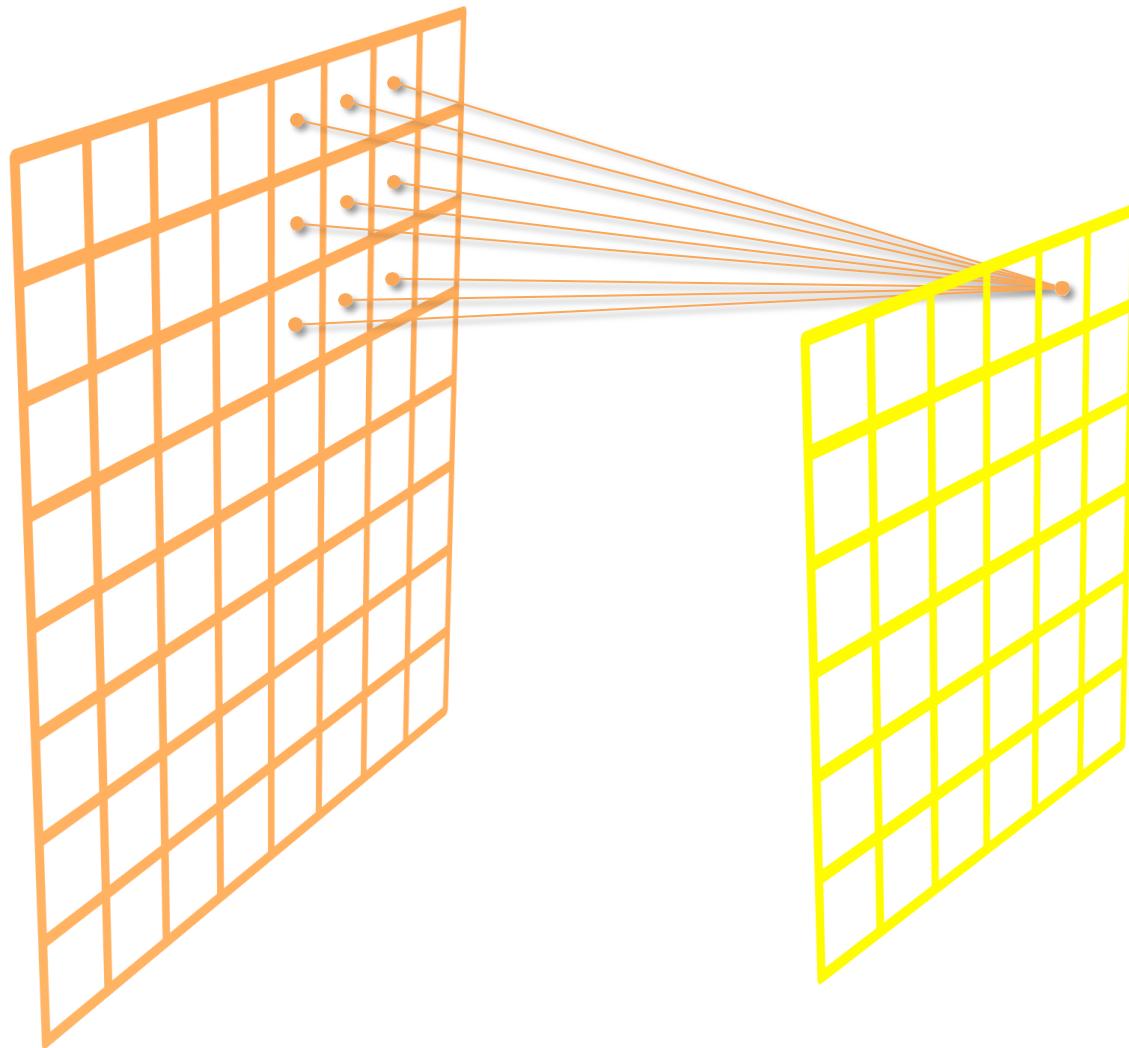
8x8 image, 3x3 filter, Stride 1



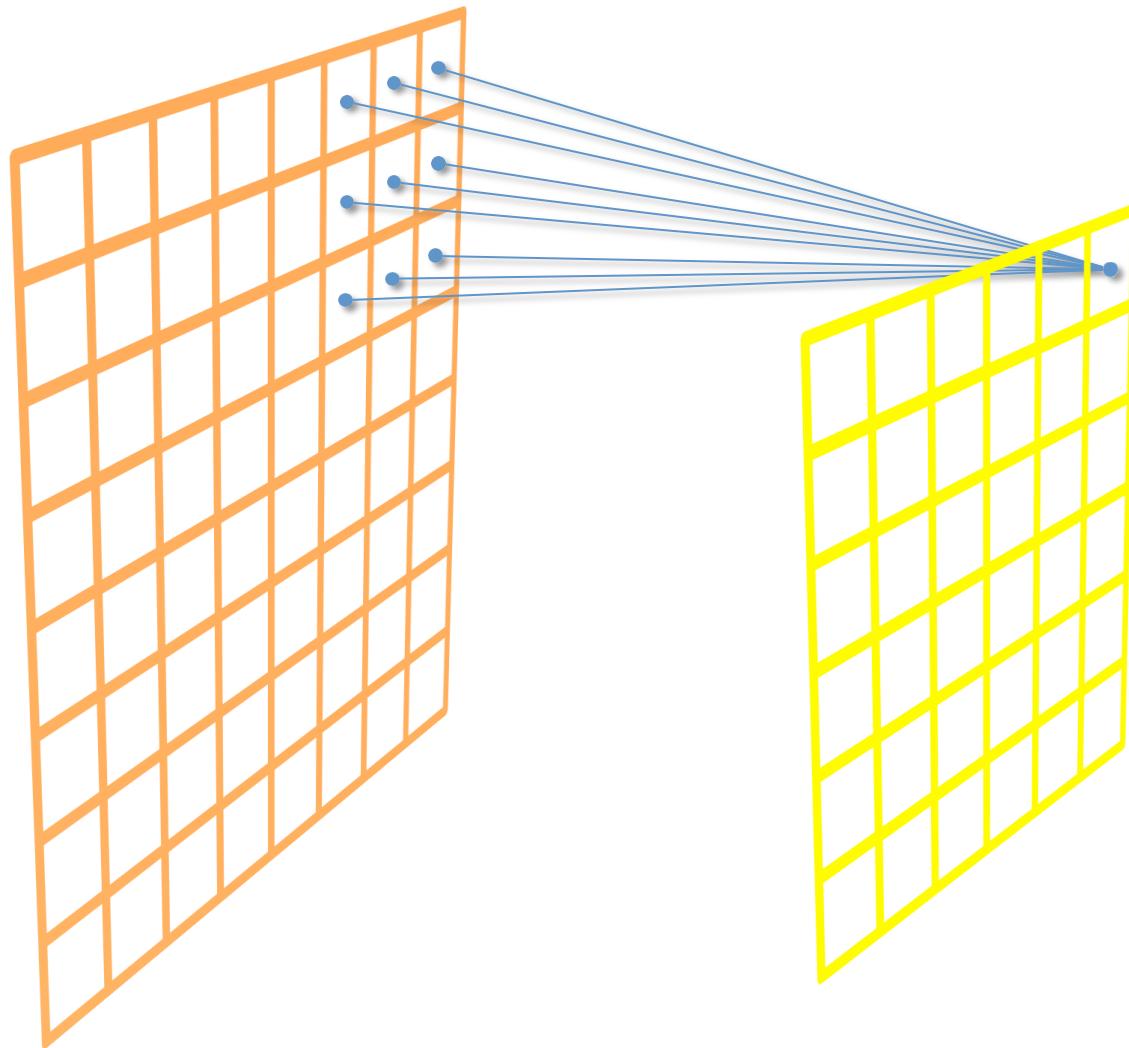
8x8 image, 3x3 filter, Stride 1



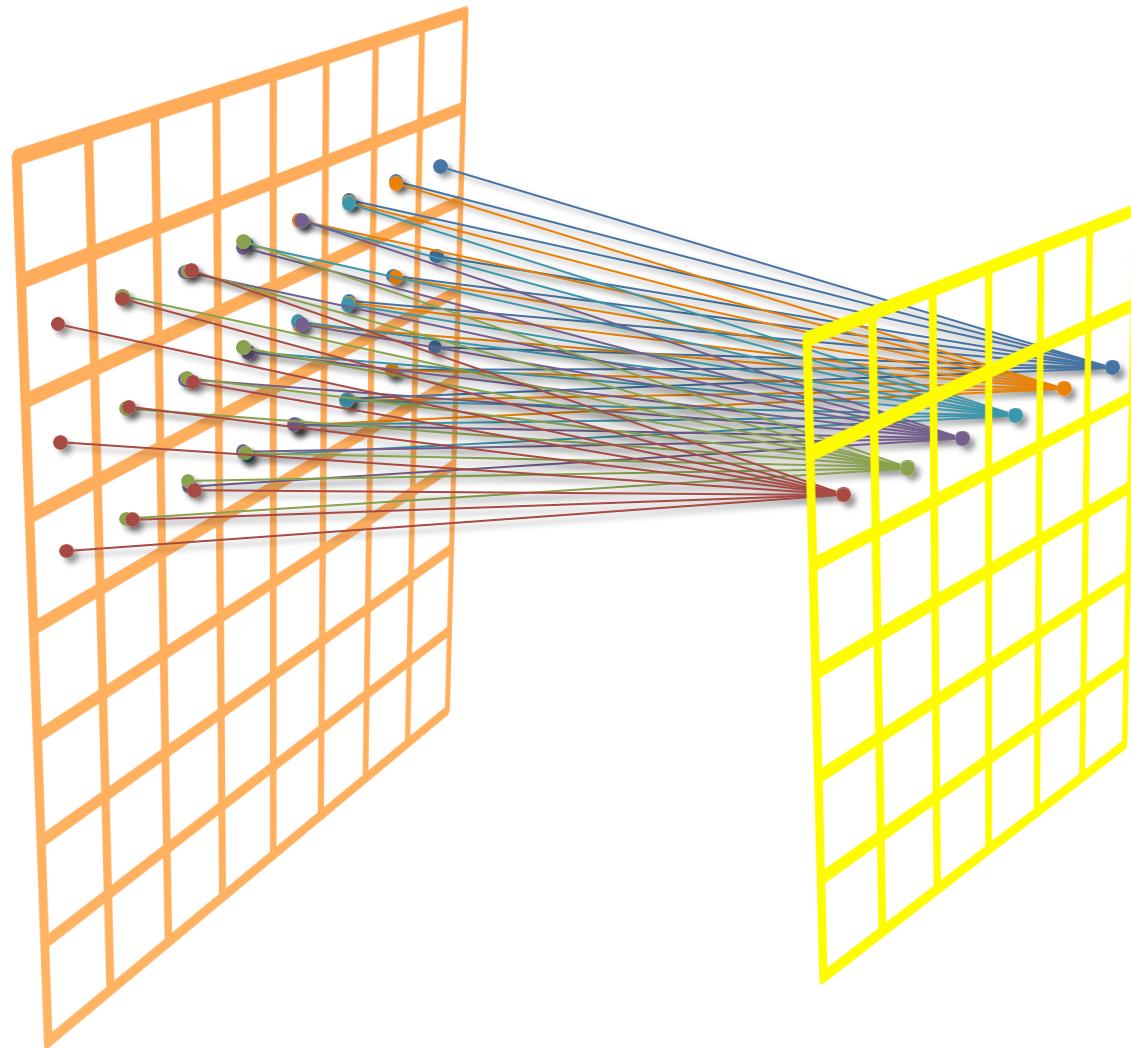
8x8 image, 3x3 filter, Stride 1



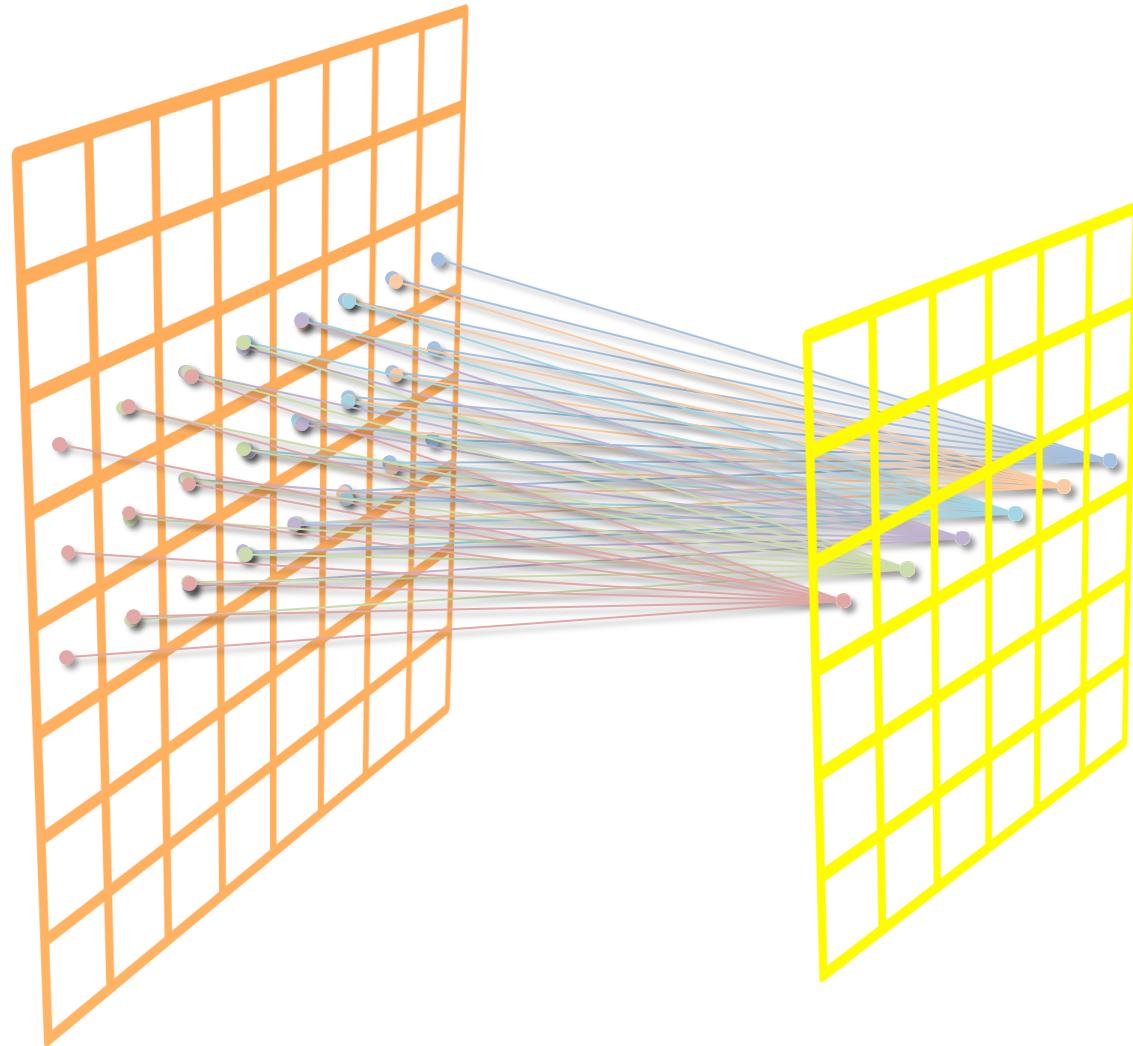
8x8 image, 3x3 filter, Stride 1

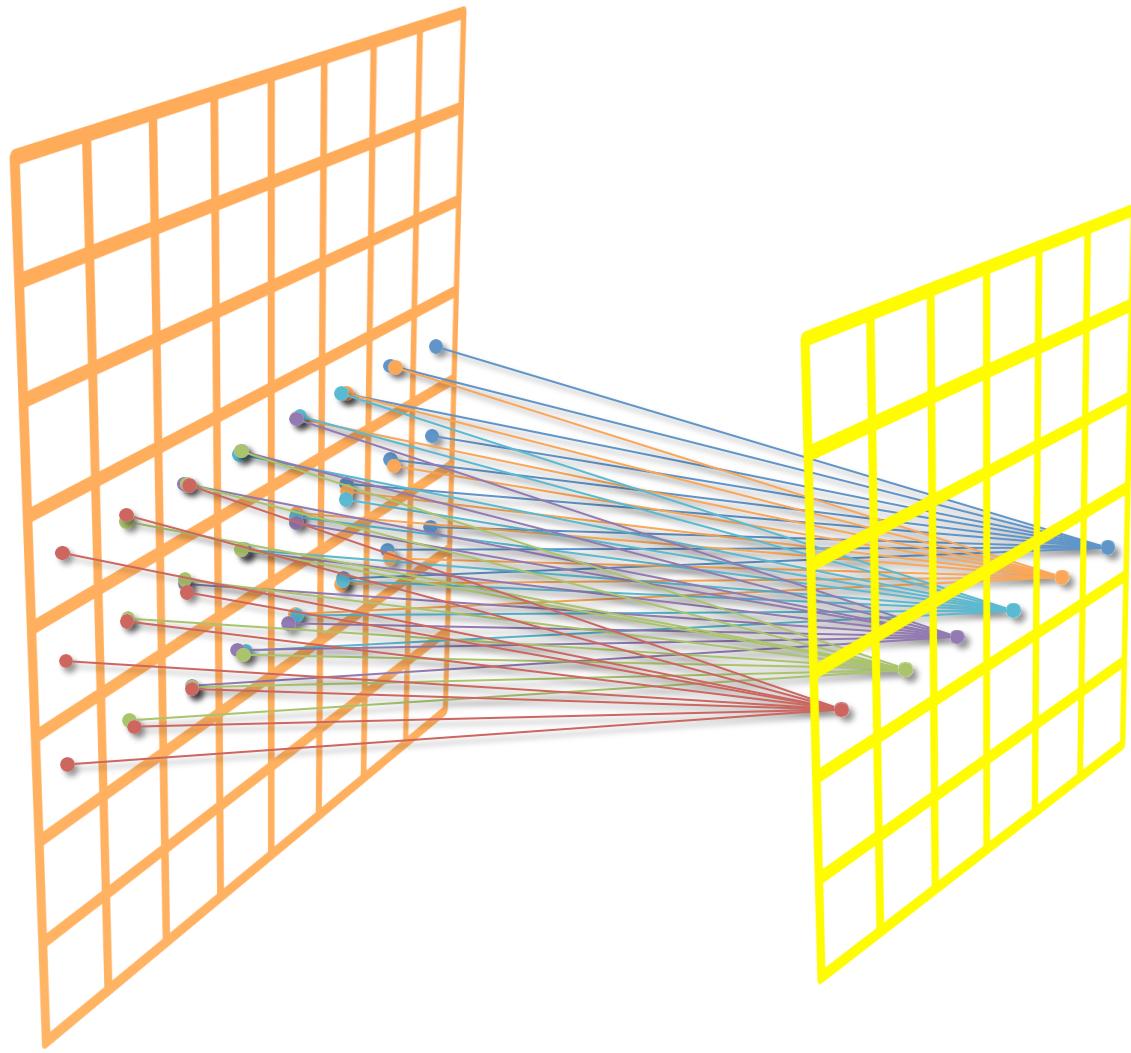


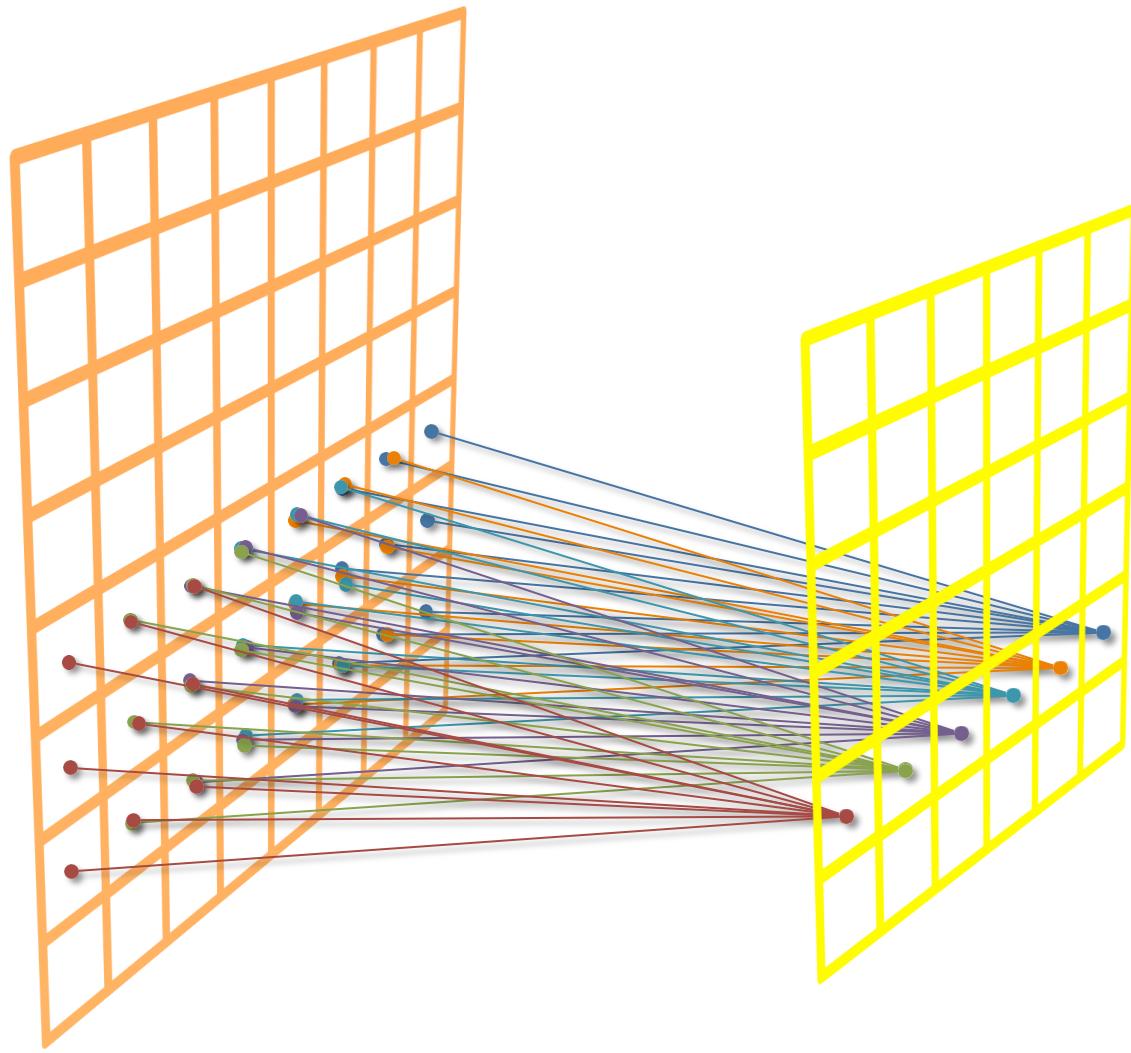
8x8 image, 3x3 filter, Stride 1



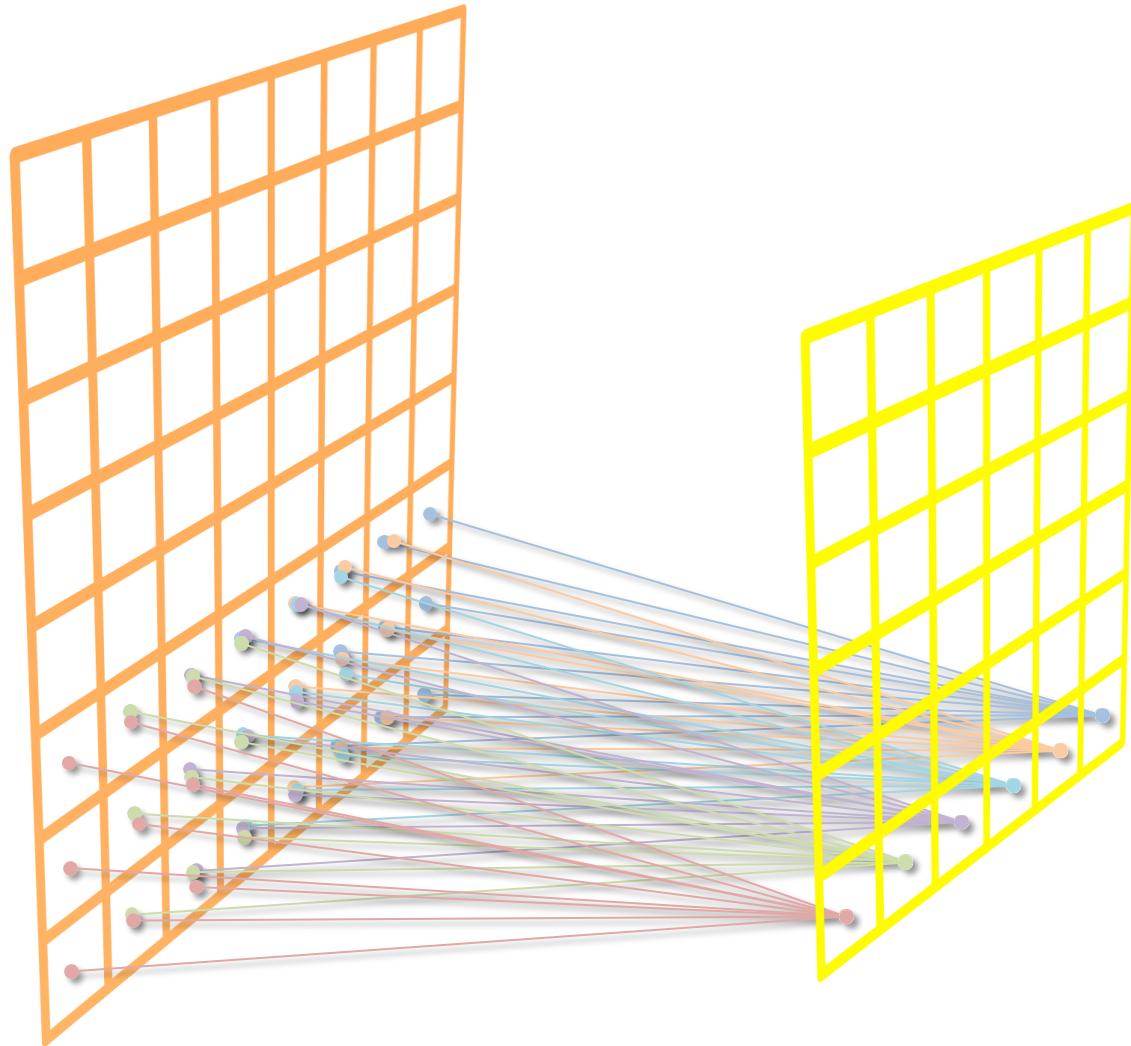
8x8 image, 3x3 filter, Stride 1



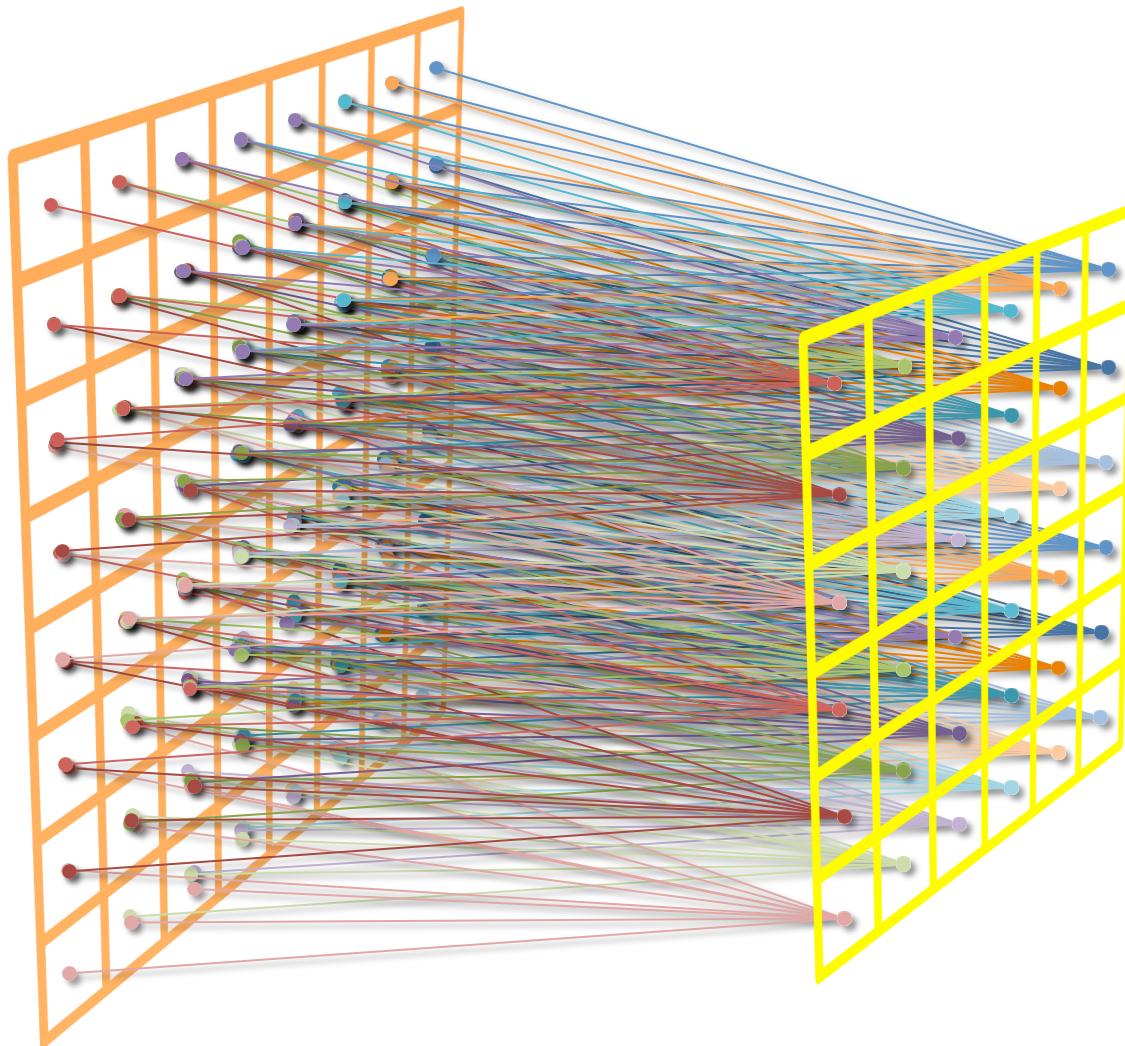




8x8 image, 3x3 filter, Stride 1



8x8 image, 3x3 filter, Stride 1



Total Connections?

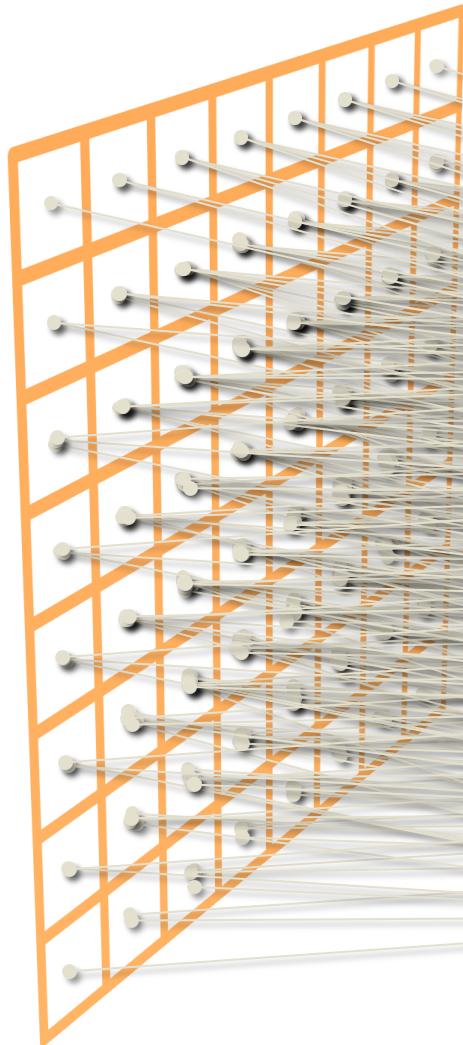
$$6 \times 6 \times 1 \quad \times \quad 3 \times 3 \times 1$$

What would have been the number of connections if this had been a fully-connected layer?

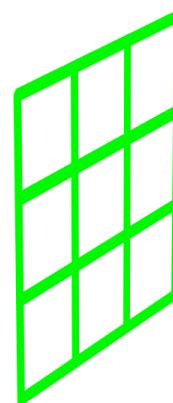
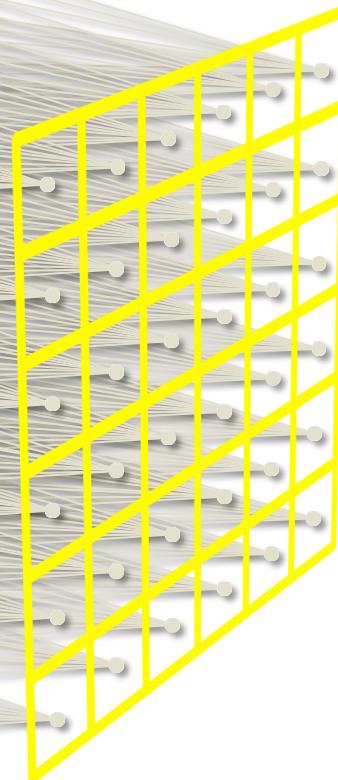
Total Unique Parameters?

$$3 \times 3 \times 1 + \text{bias}$$

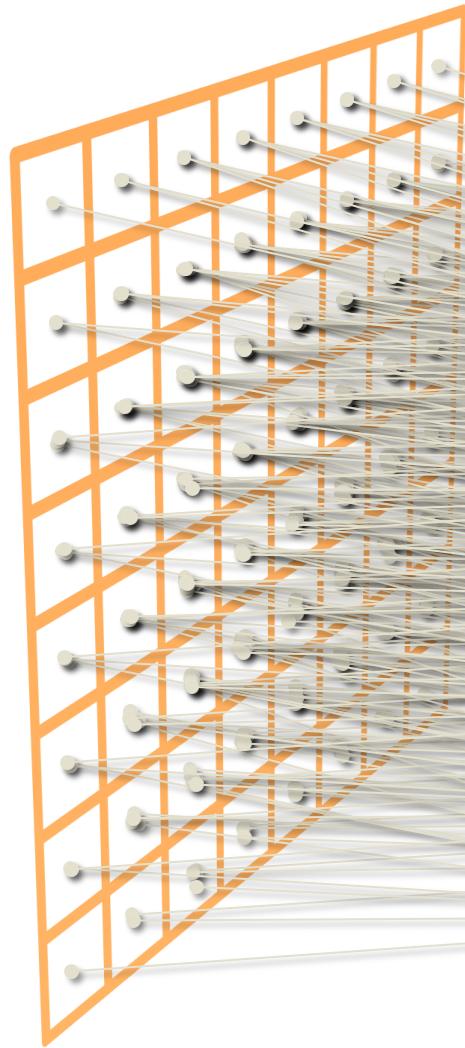
8x8 image, 3x3 filter, Stride 1



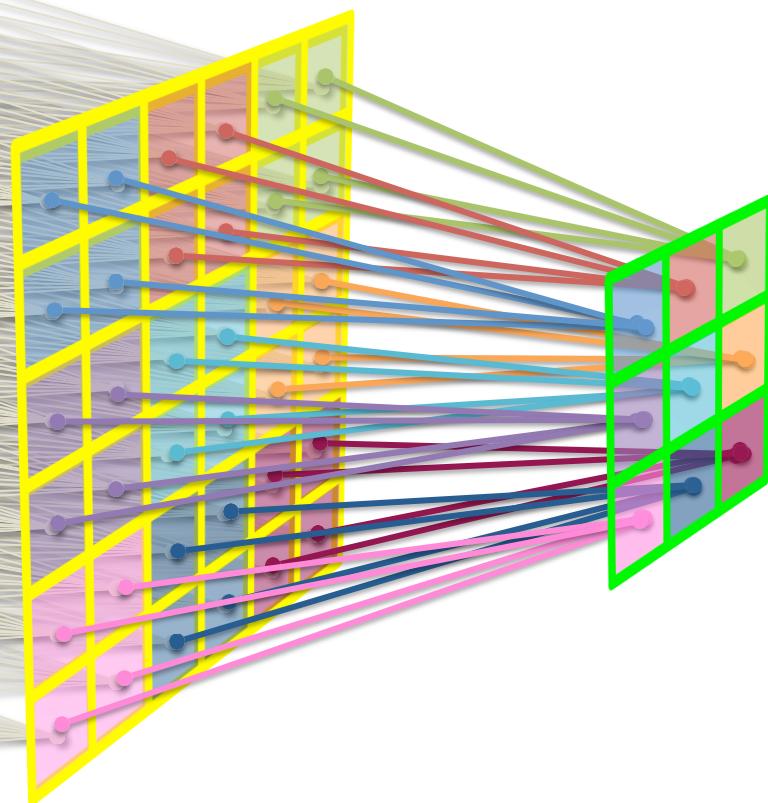
2x2 pooling, Stride 2



8x8 image, 3x3 filter, Stride 1



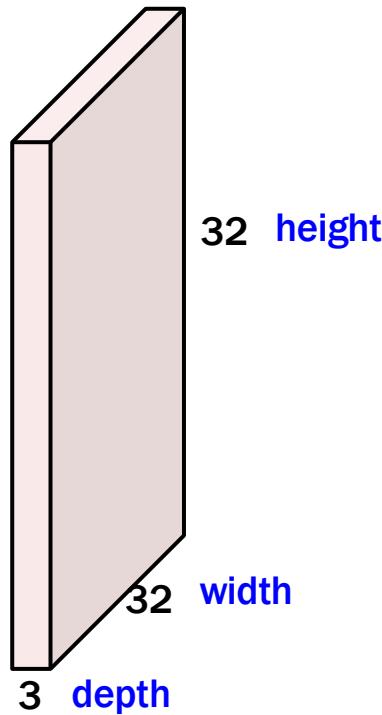
2x2 pooling, Stride 2



How Color Channels are Handled?

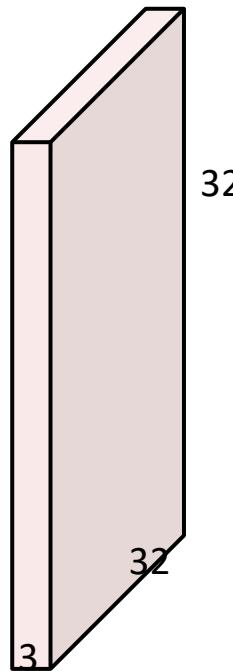
Convolution Layer

32x32x3 image



Convolution Layer

32x32x3 image

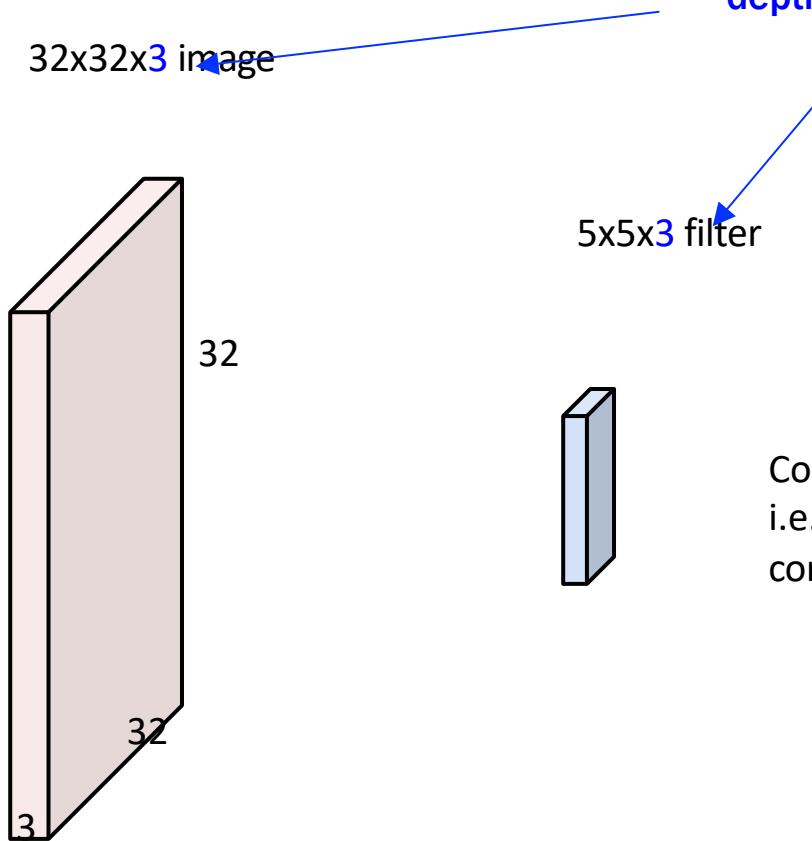


5x5x3 filter



Convolve the filter with the image
i.e. “slide over the image spatially,
computing dot products”

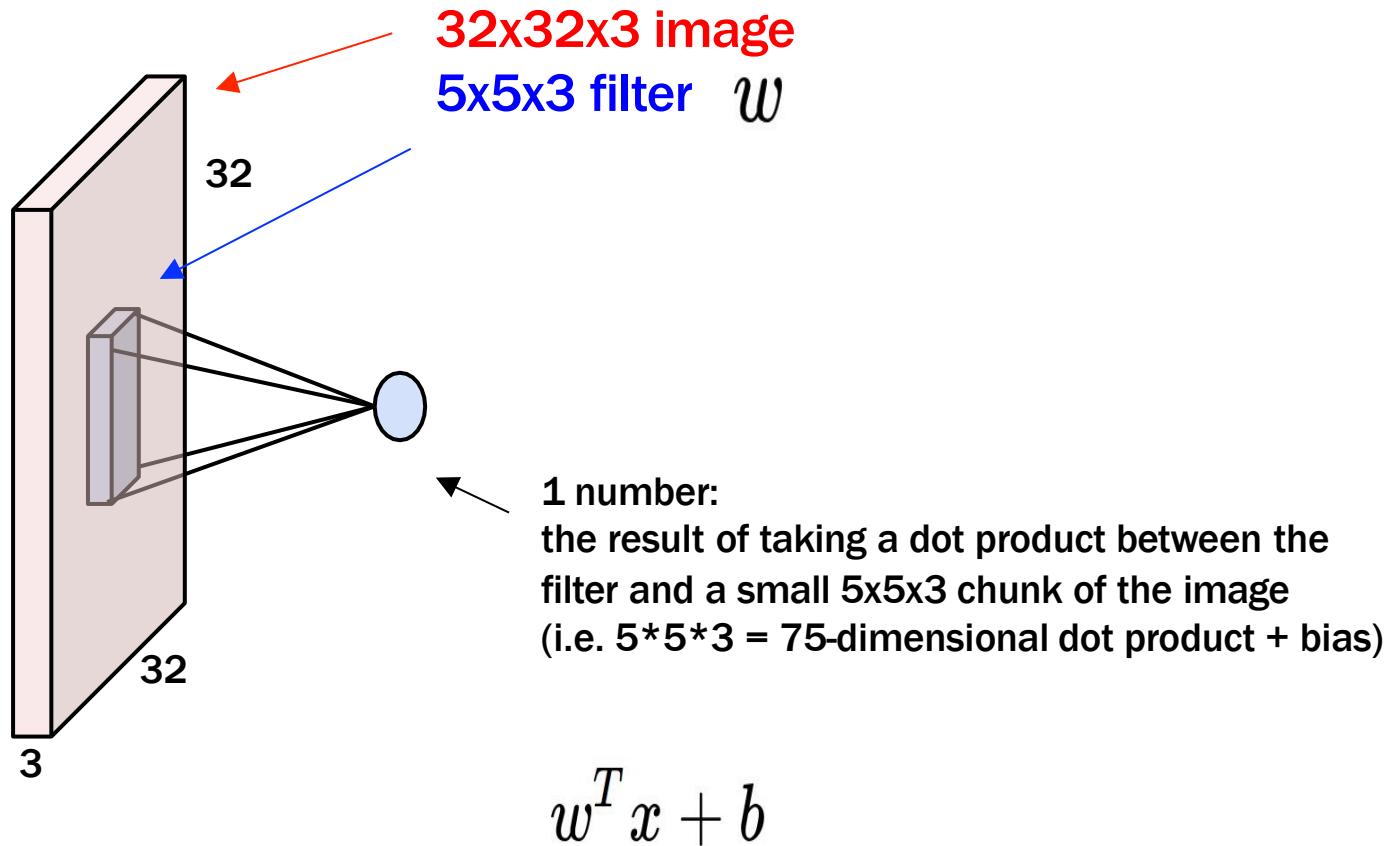
Convolution Layer



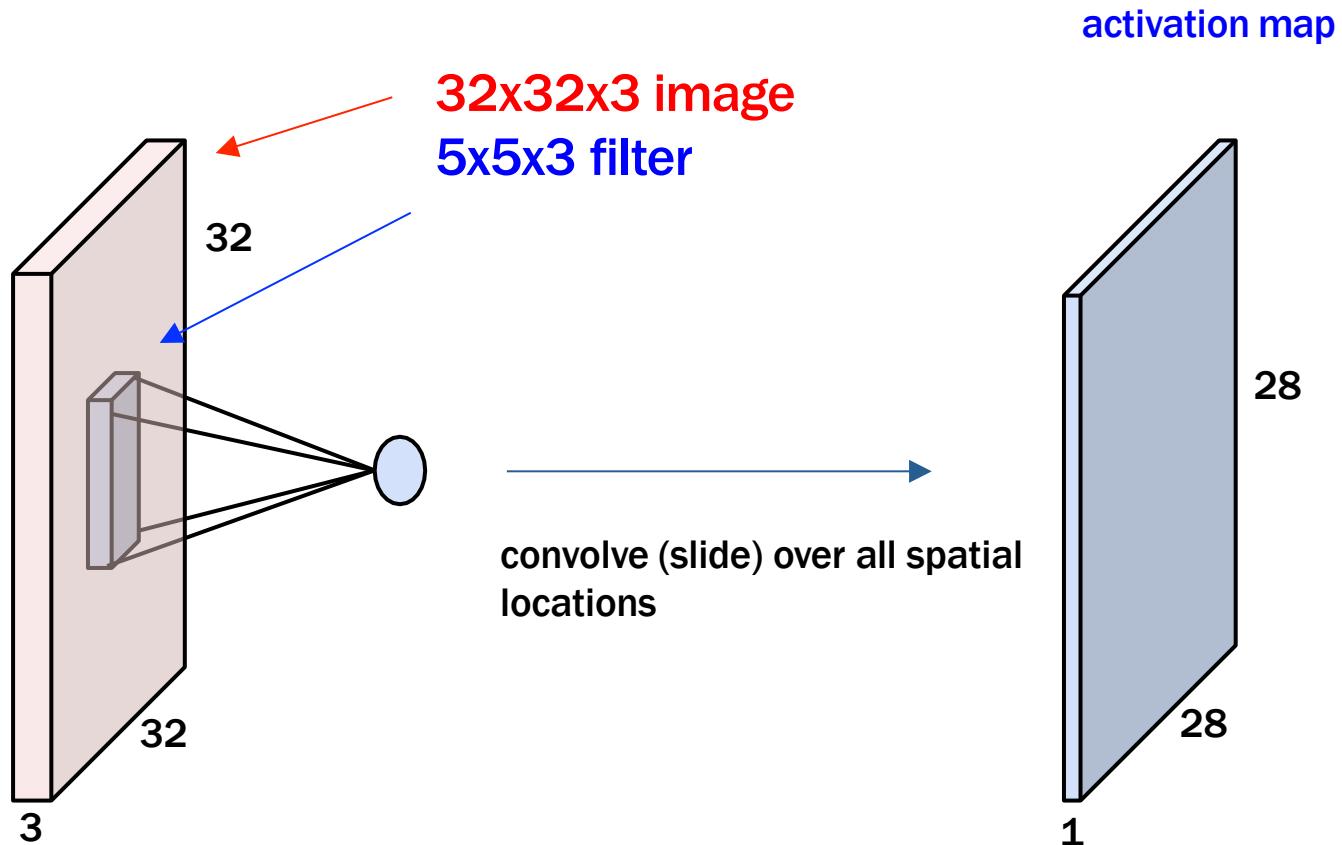
Filters always extend the full depth of the input volume

Convolve the filter with the image
i.e. “slide over the image spatially,
computing dot products”

Convolution Layer

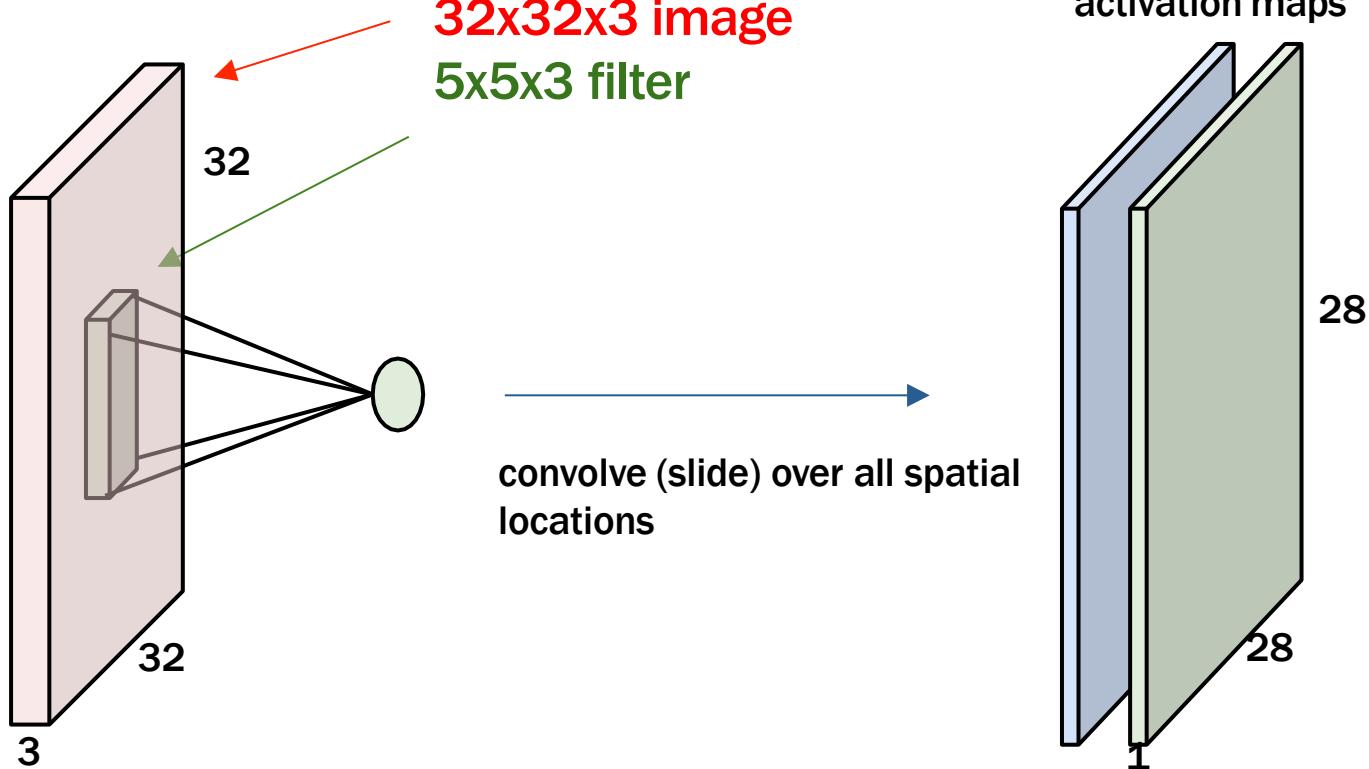


Convolution Layer

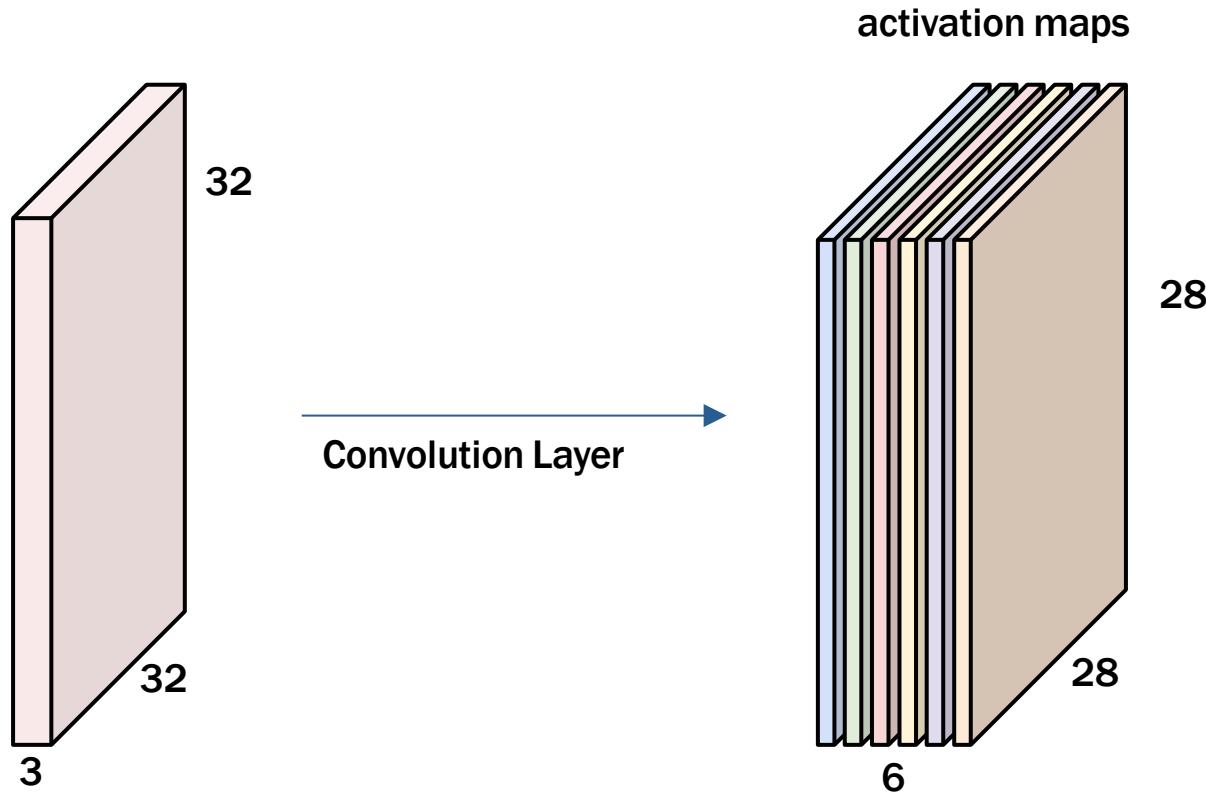


Convolution Layer

consider a second, green filter

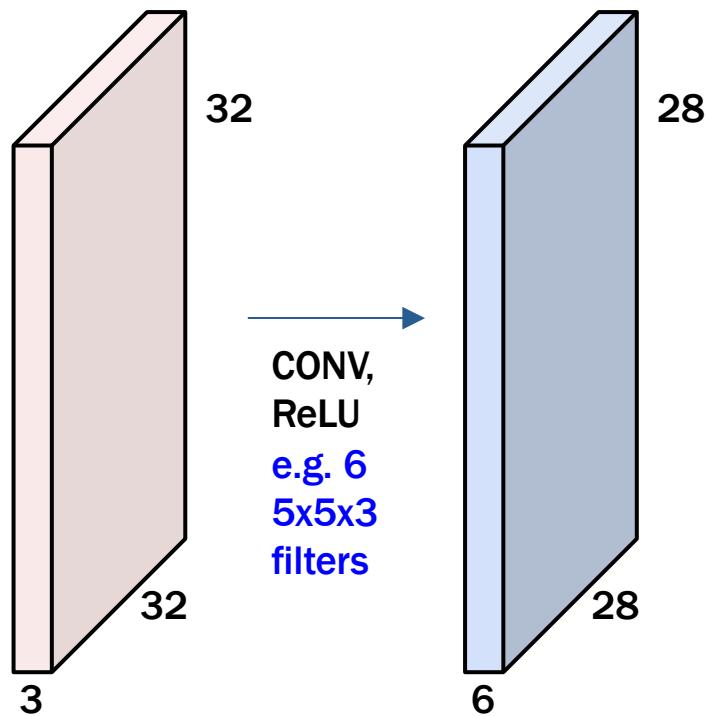


For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:

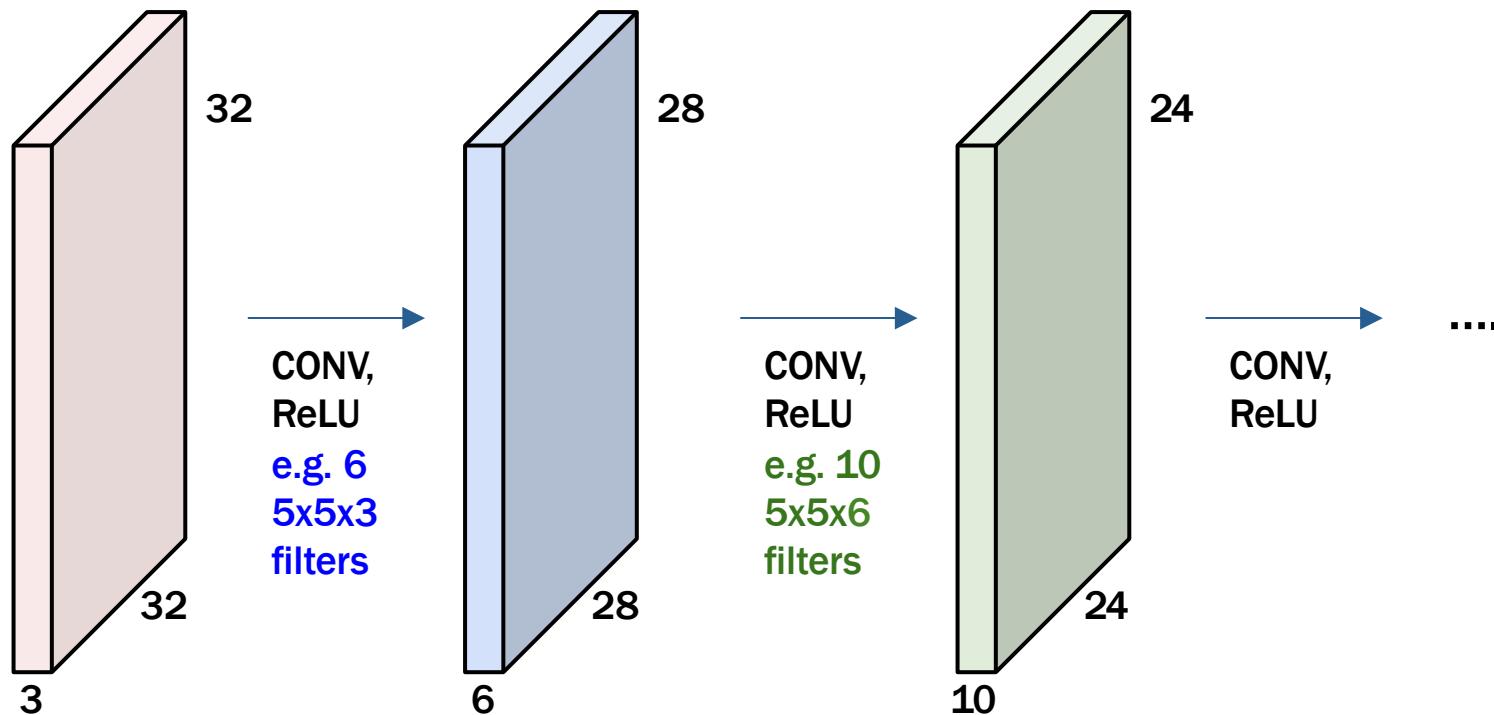


We stack these up to get a “new image” of size $28 \times 28 \times 6$!

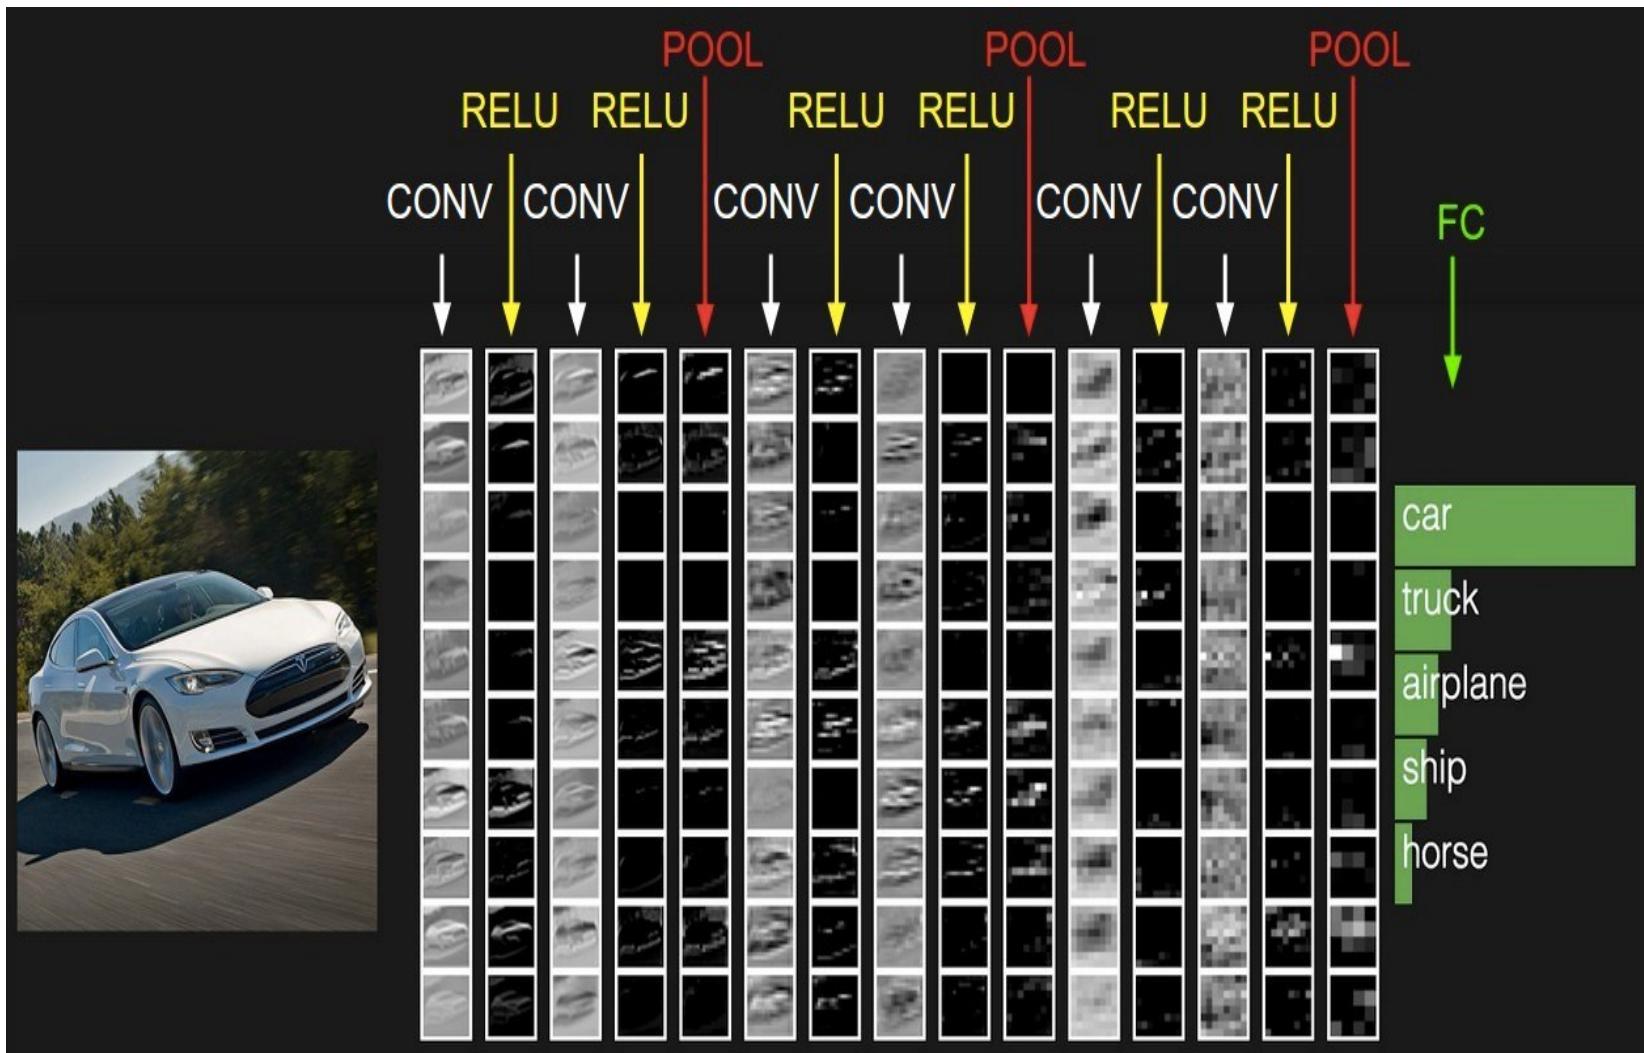
Preview: ConvNet is a sequence of Convolution Layers, interspersed with activation functions



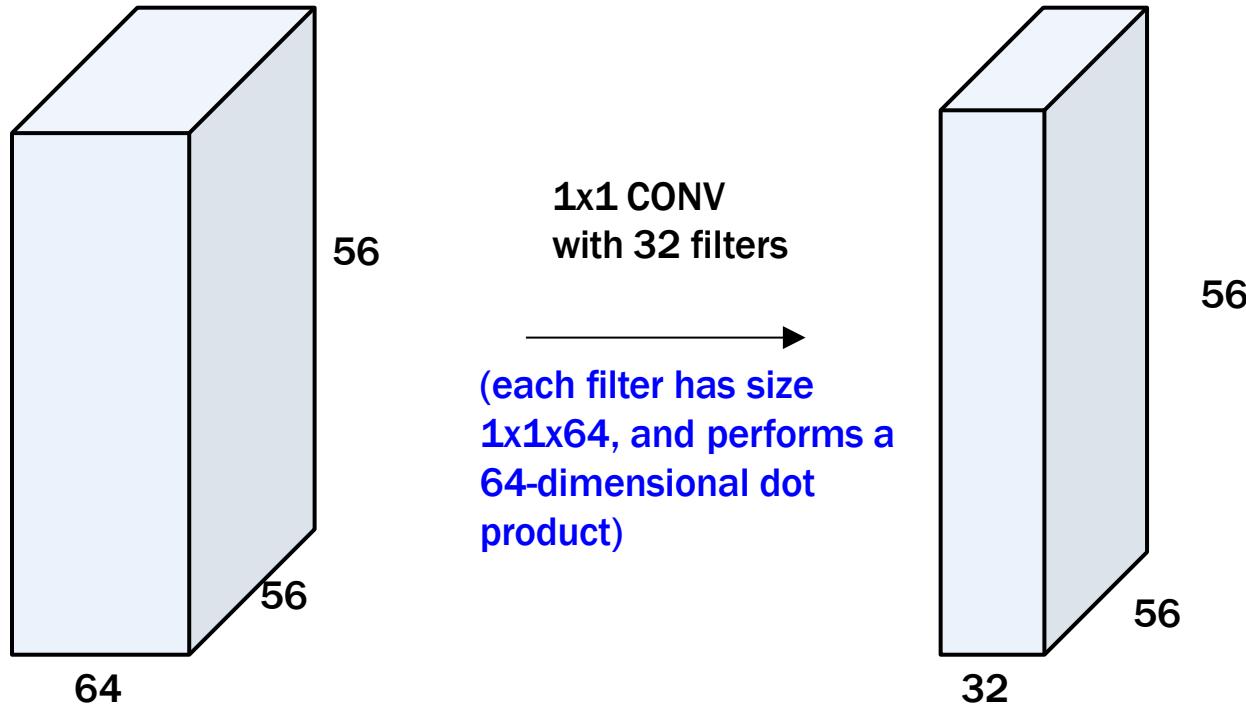
Preview: ConvNet is a sequence of Convolutional Layers, interspersed with activation functions



preview:



(btw, 1x1 convolution layers make perfect sense)



Demo

Example predictions on Test set

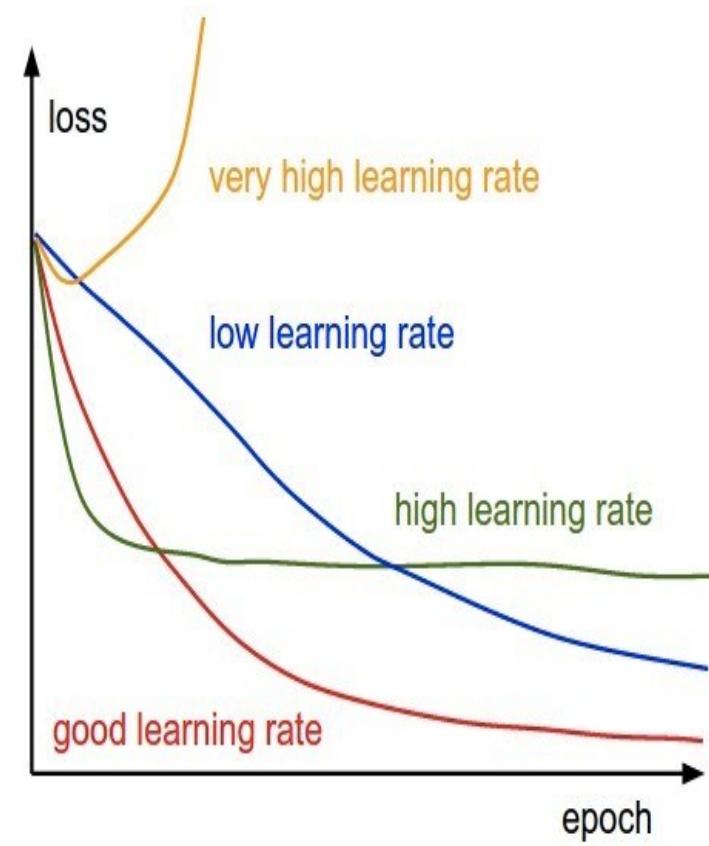
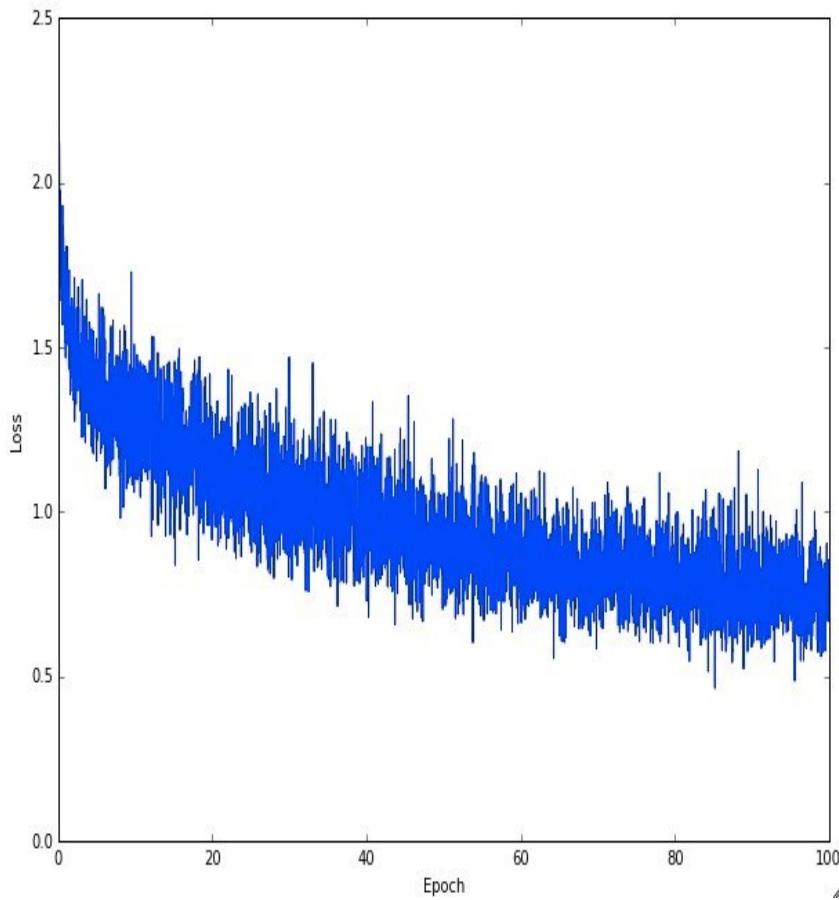
test accuracy based on last 200 test images: 0.45

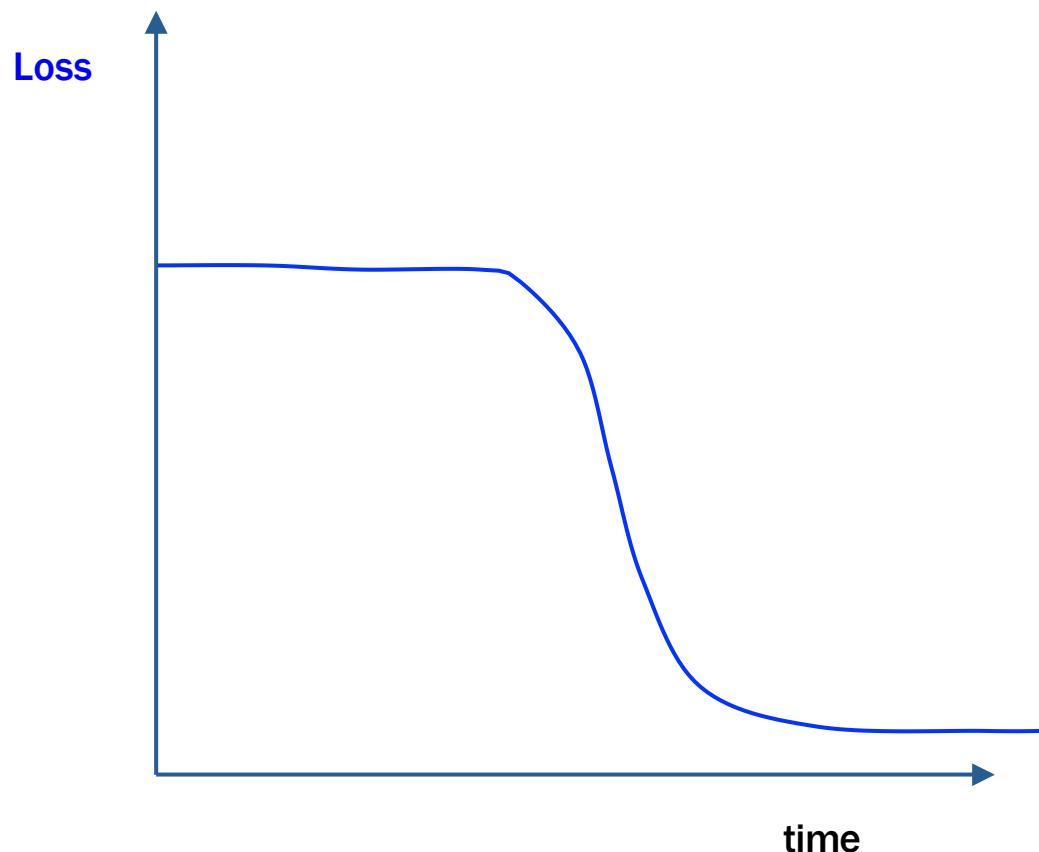
 car ship airplane	 car truck ship	 ship airplane car	 deer airplane bird
 deer cat horse	 ship airplane horse	 horse deer bird	 cat truck horse
 deer bird frog	 airplane deer bird	 deer dog cat	 airplane bird deer
 frog deer cat	 dog cat ship	 bird deer frog	 frog bird deer

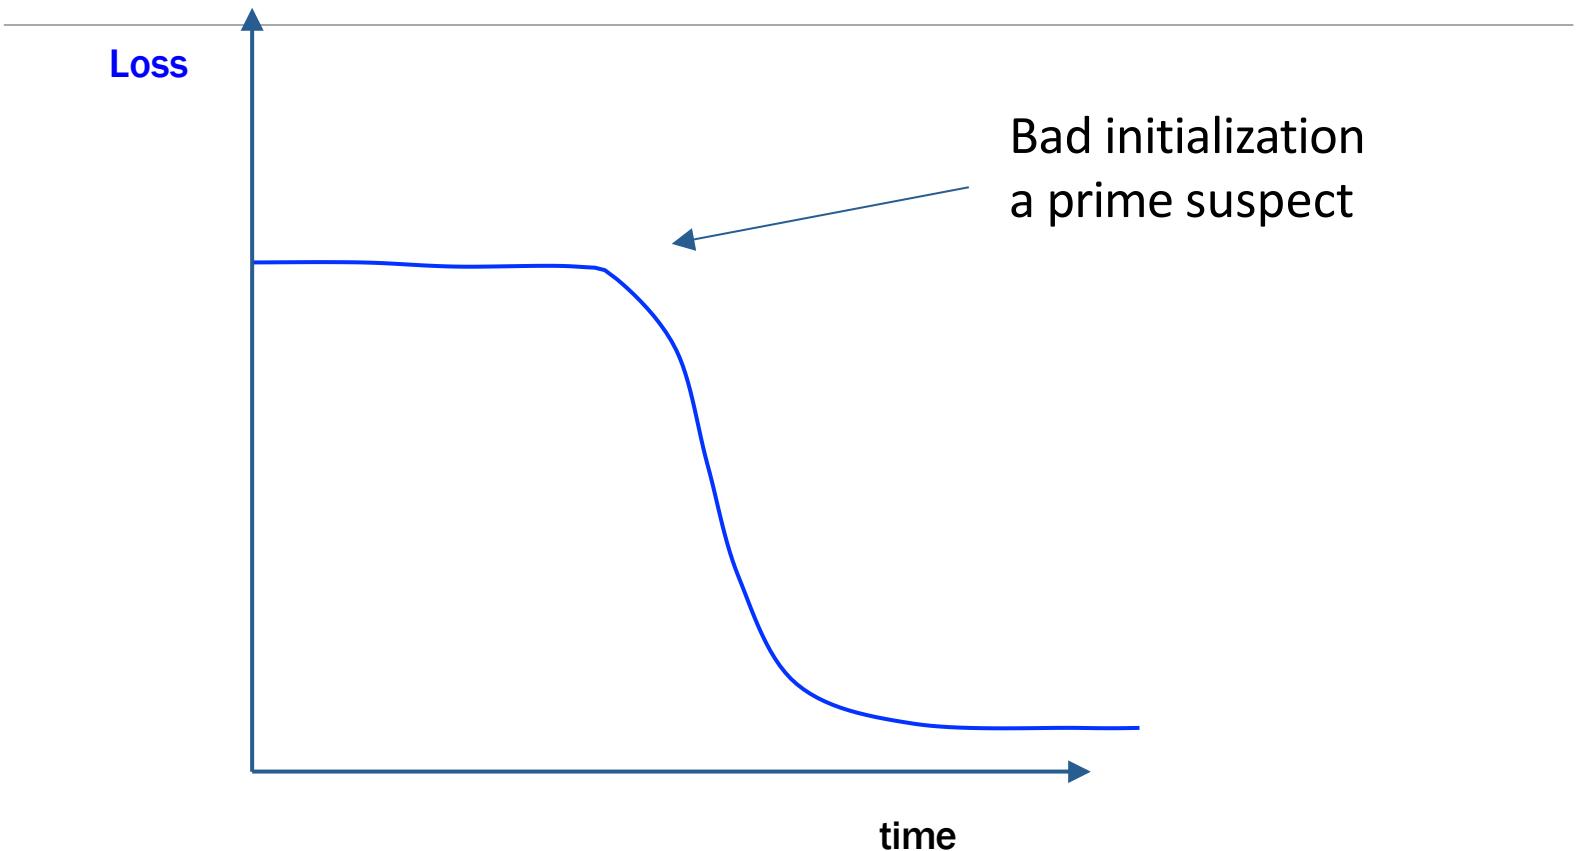
<http://cs.stanford.edu/people/karpathy/convnetjs/demo/cifar10.html>

How to Tune Hyperparameters?

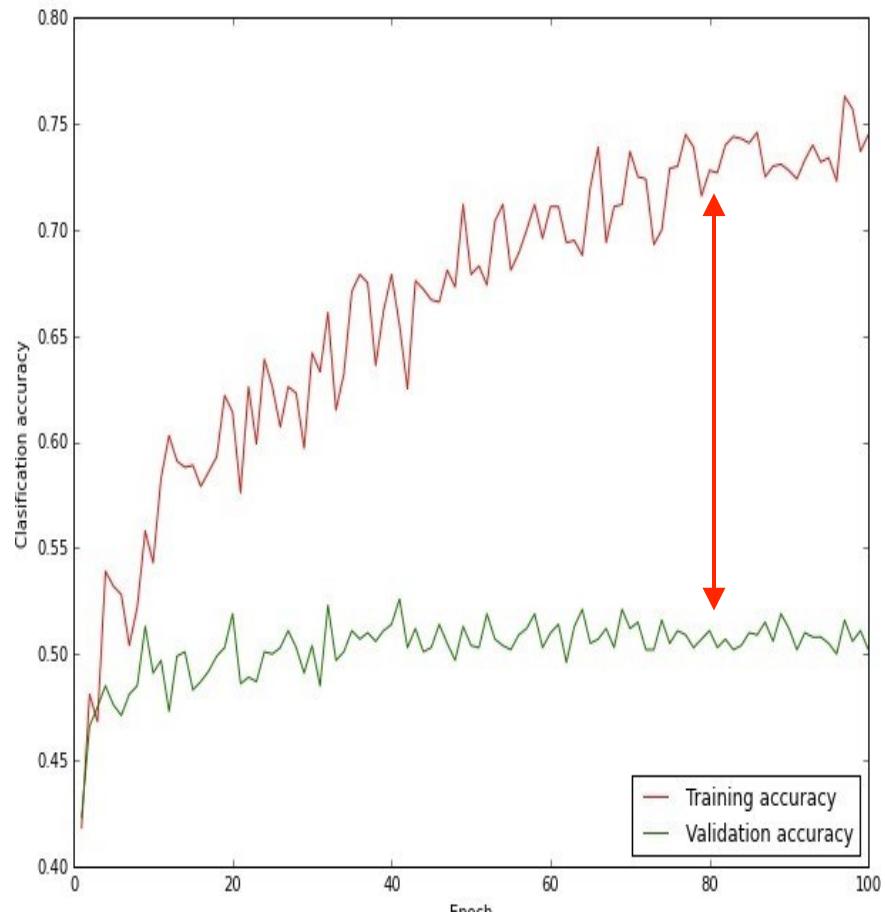
Monitor and visualize the loss curve







Monitor and visualize the accuracy:



big gap = overfitting
=> increase regularization strength?

no gap
=> increase model capacity?

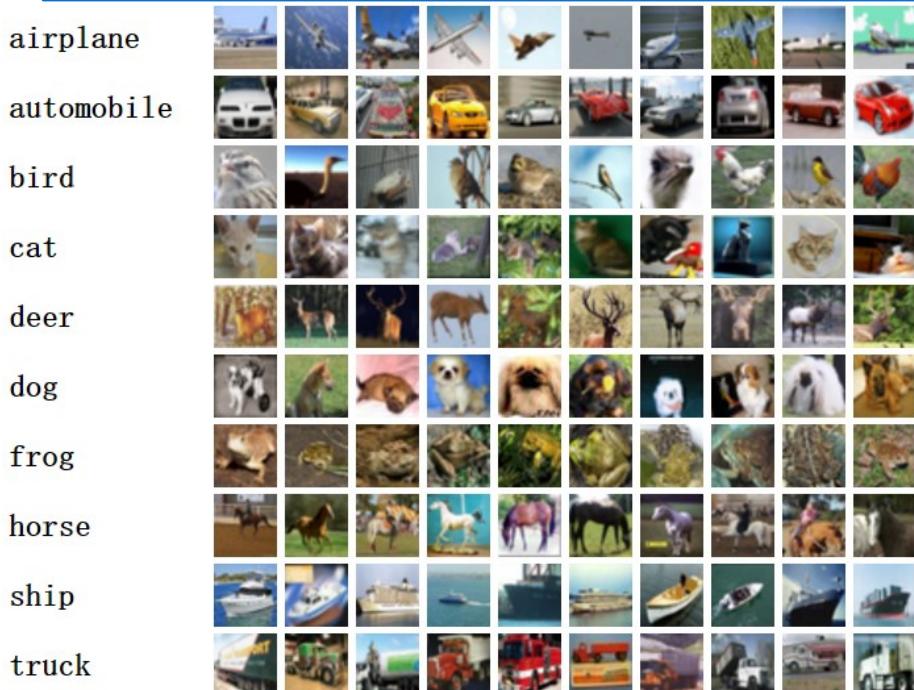
Pretrained Models

ImageNet Dataset

- ❑ The ImageNet project is a large visual database designed for use in visual object recognition software research.
- ❑ As of 2016, **over ten million URLs of images have been hand-annotated** by ImageNet to indicate what objects are pictured; in at least one million of the images, **bounding boxes** are also provided.
- ❑ The database of annotations of third-party image URL's is freely available directly from ImageNet; however, the actual images are not owned by ImageNet.
- ❑ Since 2010, the ImageNet project runs an annual software contest, the **ImageNet Large Scale Visual Recognition Challenge (ILSVRC)**, where software programs compete to correctly classify and detect objects and scenes.

CIFAR10 dataset and Results of Different Models

The CIFAR-10 dataset consists of 60,000 32x32 color images in 10 classes, with 6000 images per class. There are 50000 training images and 10,000 test images.



Accuracy

Model	Acc.
VGG16	92.64%
ResNet18	93.02%
ResNet50	93.62%
ResNet101	93.75%
ResNeXt29(32x4d)	94.73%
ResNeXt29(2x64d)	94.82%
DenseNet121	95.04%
PreActResNet18	95.11%
DPN92	95.16%

Pretrained Models

- ❑ Many pretrained models exist on ImageNet dataset
 - https://www.tensorflow.org/api_docs/python/tf/keras/applications

- ❑ **How those are useful for us?**

Book Reading

- ❑ Murphy – Chapter 8
- ❑ Jurafsky – Chapter 5, Chapter 4, Chapter 7
- ❑ Tom Mitchel – Chapter 4